

Parallel programming

Paolo Burgio
paolo.burgio@unimore.it

I still have no idea



what I'm doing



Definitions

Parallel computing

- › Cooperative
- › Partition computation across different compute engines
- › E.g., PThreads w/Shared mem, but also multi-process on the same machine!

Distributed computing

- › Competitive
- › Partition computation across different machines
- › E.g., multiprocess (MPI, MQTT) w/message passing



Why do we need parallel computing?

Increase performance of our machines

› Scale-up

- Solve a "bigger" problem in the same time
 - Metric is **throughput**

› Scale-out (ex: Speedup)

- Solve the same problem «in less time» (or, less power..)
- Metric is **end-to-end latency**



Yes but..

Why (highly) parallel machines...

...and not faster single-core machines?

The answer #1 - Money





The answer #2 – the "hot" one

Moore's law

- › "The number of transistors that we can pack in a given die area doubles every 18 months"

Dennard's scaling

- › "Performance per watt of computing is growing exponentially at roughly the same rate"


The answer #2 – the "hot" one

- › SoC design paradigm




- › Gordon Moore
 - His law is still valid, but...

- › *"The free lunch is over"*



In other words...





Instead of going faster..

- › ..(go faster but through) parallelism!

Problem #1

- › New computer architectures
- › At least, three architectural templates

Problem #2

- › Need to efficiently program them
- › HPC already has this problem!

The problem

- › Programmers **must** know a bit of the architecture!
- › To make parallelization effective
- › "Let's run this on a GPU. It certainly goes faster" (cit.)

The Big problem

- › Effectively programming in parallel is difficult

Brian Kernighan (1942-)

- Researcher, theory of informatics
- Co-authored *UNIX* and *AWK*
- Wrote "*The C Programming Language*" book

“Everyone knows that debugging is twice as hard as writing a program in the first place.

So if you’re as clever as you can be when you write it, how will you ever debug it?”



Amdahl's law

- › A sequential program that takes 100 sec to exec
- › Only 95% can run in parallel (it's **a lot**)
- › And.. you are an **extremely good** programmer, and you have a machine with **1billion cores**, so that part takes **0 sec**
- › So,

$$T_{par} = 100_{sec} - 95_{sec} = 5_{sec}$$

$$Speedup = \frac{100_{sec}}{5_{sec}} = 20x$$


...20x, on one billion cores!!!



Symmetric multi-processing

- › Memory: centralized with bus interconnect, I/O
- › Typically, **multi-core** (sub)systems
 - Examples: Sun Enterprise 6000, SGI Challenge, Intel (this laptop)


Can be 1 bus, N
busses, or any
network



Asymmetric multi-processing


- › Memory: centralized with uniform access time (UMA) and bus interconnect, I/O
- › Typically, multi-core (sub)systems
 - Examples: ARM Big.LITTLE, NVIDIA Tegra X2 (Drive PX)

Can be 1 bus, N
busses, or any
network



SMP – distributed shared memory

- › Non-Uniform Access Time - **NUMA**
- › **Scalable** interconnect
 - Typically, **many** cores
 - Examples: embedded accelerators, GPUs




UMA vs. NUMA


- > Shared
- (
- > Uniform
- D

	MEM0	MEM1	MEM2	MEM3
CPU0...3	0 clock	10 clock	20 clock	10 clock
CPU4...7	10 clock	0 clock	10 clock	20 clock
CPU8...11	20 clock	10 clock	0 clock	10 clock
CPU12..15	10 clock	20 clock	10 clock	0 clock

UMA



NUMA



Programming abstractions





What is...

- › ...a core?
- › ...a program?
- › ...a process?
- › ...a thread?
- › ...a task?





What is...

- › ...a core?
 - An electronic circuit to execute instruction (=> programs)
- › ...a program?
 - The implementation of an algorithm
- › ...a process?
 - A program that is executing
- › ...a thread?
 - A unit of execution (of a process)
- › ...a task?
 - A unit of work (of a program)



What is...

- › ...a core?
 - An electronic circuit to execute instruction (=> programs)
- › ...a program?
 - The implementation of an algorithm
- › ...a process?
 - A program that is executing
- › ...a thread?
 - A unit of execution (of a process)
- › ...a task?
 - A unit of work (of a program)






(Simplest) threading model

Fork-join execution model


- › The main, single thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel (as a team of threads)
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // >>> PARALLEL REGION
                             &myattr,
                             my_pthread,
                             NULL);
    // PARALLEL REGION.
    // Here, the main thread can do other stuff

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &ret); // <== JOIN
}
```



Let's see
this in
action




Master-slave threading model

Fork-join execution model

- › The main, **Master** thread spawns a team of **Slave** threads (here, NTHREADS = 3)
- › They all perform computation in parallel (as a team of threads)
- › At the end, they are joined one by one (aka: barrier)

```
int main()
{
    int err;
    pthread_t mythreads[NTHREADS];
    for (int i=0; i<NTHREADS; i++)
        err = pthread_create (&mythreads[i], // ==> FORK
                            &myattr,
                            my_pthread_fn,
                            NULL);
    // PARALLEL REGION.
    // Here, the main thread can do other stuff!
    other_fn(); // or can run my_pthread_fn

    for (int i=0; i<NTHREADS; i++)
        pthread_join(mythreads[i], &returnvalue); // <== JOIN
}
```




Work partitioning

Several models (aka: patterns), here to cite a few

- › Data parallelism (see also GPGPUS)
 - We're getting there, don't worry...
 - Reduction
- › Task parallelism
 - Work queue
 - Pipelining
- › Offloading

...and a mix of them...



Data parallelism

(Aka: data decomposition, loop decomposition, SPMD, SIMD*...)

Parallel threads execute the same operation(s) on multiple data

- › Data is typically an array, a matrix (image)....
 - Note: you typically map the iteration id of the loop to the data index
- › **Partitioning strategy** defines how many iterations (chunk) every thread will perform
 - From 1 iteration, to loop size



* Single Program, Multiple Data; Single Instruction, Multiple Data



Exercise

Let's
code!

Create an array of N elements

- › Put inside each array element its index, multiplied by '2'
- › `arr[0] = 0; arr[1] = 2; arr[2] = 4; ...and so on..`

Now, do it in parallel with a team of T threads


- › Assume N is always bigger than T
- › "Decompose" the `for` construct, so that every thread manages (chunk size is) N/T iterations



Loop partitioning among threads

Let's code!

- › Case #1: N multiple of T
 - Say, N = 20, T = 4
- › chunk = #iterations for each thread
- › Very simple..




$$chunk = \frac{N}{T} ;$$

$$i_{start} = thread_{ID} * chunk; \quad i_{end} = i_{start} + chunk$$

Loop partitioning among threads

Let's code!

- › Case #2: N **not** multiple of T
 - Say, N = 19, T = 4
- › chunk = #iterations for each thread (but last)
 - Last thread has less! ($chunk_{last}$)



$$chunk = \frac{N}{T} + 1; \quad chunk_{last} = N \% chunk$$

$$i_{start} = thread_{ID} * chunk; \quad i_{end} = \begin{cases} i_{start} + chunk & \text{if not last thread} \\ i_{start} + chunk_{last} & \text{if last thread} \end{cases}$$



"Last thread"

Let's
code!

- › Unfortunately, we don't know which thread will be "last" in time
- › But...we don't actually care the order in which iterations are executed
 - If there are not dependencies..
 - And..we do know that


$0 \leq \text{myid} < \text{NUM_THREADS}$

- › We **choose** that last thread as highest number

Loop partitioning among threads

Let's
code!

- › Case #3: N **smaller than T**
 - Say, N = 3, T = 4
- › Similar to Case #1
 - With some adjustments...





Let's put them together!

Let's
code!

- › Case #1 (N multiple of T)

$$chunk = \frac{N}{T} \quad i_{start} = thread_{ID} * chunk; \quad i_{end} = i_{start} + chunk$$

- › Case #2 (N not multiple of T)

$$chunk = \frac{N}{T} + 1; \quad chunk_{last} = N \% chunk$$

$$i_{start} = thread_{ID} * chunk; \quad i_{end} = \begin{cases} i_{start} + chunk & \text{if not last thread} \\ i_{start} + chunk_{last} & \text{if last thread} \end{cases}$$

Reduction

wikipedia

*The reduction clause can be used to perform some forms of **recurrence** calculations (involving **mathematically associative and commutative operators**) in parallel. For parallel [...], a **private** copy of each list item is created, one for each implicit task, as if the private clause had been used. [...] The private copy is then initialized as specified above. At the end of the region for which the reduction clause was specified, the original list item is updated by **combining its original value with the final value of each of the private copies**, using the combiner of the specified reduction-identifier.*

E.g., average value of a sequence (array/vector)

- › Create a thread-local copy of a variable
- › Accumulate sums only for the assigned part of the array/vector
- › Then, sum the partial sums (and divide by size)



Exercise


Let's
code!

Create an array of N elements

- › ..or a vector
- › Initiate it randomly
- › Now, compute its average value using multi-threading and reduction paradigm


Unbalanced loop partitioning

- › So far, we assigned iterations «statically»
 - Might not be effective nor efficient



How can we manage dynamics/irregular workloads?

- › We would like something like this..




A different parallel paradigm: Tasking

Implements a **producer-consumer paradigm**

Managed by a **task queue**

- › Where units of work (**tasks**)
- › are pushed by threads (`q_push` primitive)
- › and pulled and executed by threads (`q_pop` primitive)




«What» happens «when»?

```
void t0() {  
    // Task 0  
}  
  
void t1() {  
    // Task 1 pushes t2 in the q  
    q_push(t2());  
}  
  
void t2() {  
    // Task 2  
}  
  
void *thread_fn(void * args) {  
    // Push t0 and t1  
    q_push(t0());  
    q_push(t1());  
}
```

```
void *other_thread_fn(void * args) {  
    // Pop a task (which one?)  
    t = q_pop();  
    // Execute it  
    t();  
}
```

*pseudo-code




The queue

It's a shared resource

- › Its primitives `q_push` and `q_pop` are *thread-safe*, i.e., internally, the concurrent access to shared data structure is protected by semaphores/mutexes

Typically implemented as a FIFO queue

- › ..but we can also have more complex semantics (e.g., parent-son => DAGs) among tasks
- › We can have queues that are tailored/more efficient for a specific problem/domain/algorithim!





Parallel programming patterns

Creational/threading

- › *Fork-join*
- › *Master-slave*
- › *Team of threads*

Synchronization

- › Barrier
- › Critical section/Atomic
- › Locks, semaphores, mutexes...

Work partitioning

- › *Data-parallelism*
- › *Reduction*
- › *Tasking/queue*
- › Offloading

Behavioral

- › Producer-consumer
- › Collaborative (multi-thread)
- › Competitive (multi-process)

References



Course website

- › http://hipert.unimore.it/people/paolob/pub/Industrial_Informatics/index.html

My contacts

- › palo.burgio@unimore.it
- › <http://hipert.mat.unimore.it/people/paolob/>

Resources

- › "Parallel programming" course by "a guy" @UNIMORE
 - https://hipert.unimore.it/people/paolob/pub/Calcolo_Parallelo/index.html
 - <https://github.com/HiPeRT/cp19/>
- › Gamma, et.al «Design Patterns – Elements of reusable Object Oriented Software», Addison Wesley
- › Douglass – «Design Patterns for Embedded Systems in C», Newnes
- › A "small blog"
 - <http://www.google.com>