

# Concurrency

---

Paolo Burgio  
paolo.burgio@unimore.it



**UNIMORE**  
UNIVERSITÀ DEGLI STUDI DI  
MODENA E REGGIO EMILIA

High Performance  
Real Time **Lab**

**FRATELLI**

**QUANDO LA MAMMA ESCE DI CASA**



**QUANDO TORNA A CASA**





# Why concurrency?

---

## Functional

- › Many users may be connected to the same system at the same time
- › Each user can have its own processes that execute concurrently with the processes of the other users
- › Perform many operations concurrently

## Performance

- › Take advantage of blocking time
- › While some thread waits for a blocking condition, another thread performs another operation
- › On a multi-core machine, independent activities can be carried out on different cores at the same time



# Competitive vs. Cooperative

---

## **Competitive** concurrency

- › Different activities compete for the resources
- › One activity does not know anything about the other
- › The OS must manage the resources so to
  - Avoid conflicts
  - Be fair

## **Cooperative** concurrency

- › Many activities cooperate to perform an operation
- › Every activity knows about the others
- › They must synchronize on particular events



# Competitive

---

Competing activities need to be “protected” from each other

- › Separate memory spaces, as with different processes

The **allocation** of the resource and the synchronization must be **centralized**

- › Competitive activities request for services to a central manager (the OS or some dedicated process) which allocates the resources in a fair way

Client/Server model

- › Communication is usually done through **messages**

More suitable to the **process** model of execution



# Client/server model

---

A server manages the resource **exclusively**

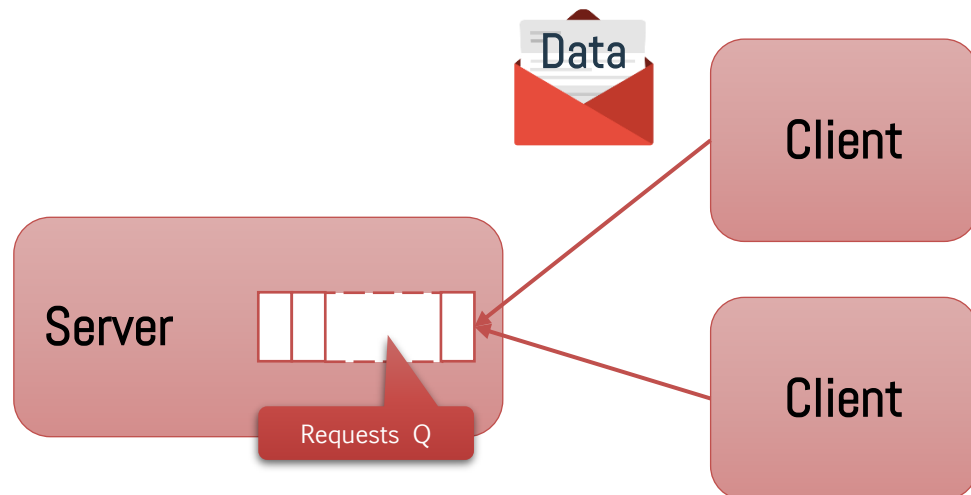
- › For example, the printer

If a process needs to access the resource, it sends a **request to the server**

- › For example, printing a file, or asking for the status
- › The server can send back the responses
- › The server can also be on a remote system

Two basic primitives:

- › **send** and **receive**





# Cooperative model

---

Cooperative activities **know** about each other

- › They do not need memory protection (less overhead)

They need to access the same data structures

- › Allocation of the resource is **de-centralized**
- › **Shared memory** model

More suitable to the **thread** model of execution



# Competition vs. cooperation

---

**Competition** is best resolved by using the **message passing** model

- › However, it can be implemented using a shared memory paradigm too

**Cooperation** is best implemented by using the **shared memory** paradigm

- › However, it can be realized by using pure message passing mechanisms

General purpose OS needs to support both models

- › Protection for competing activities
- › Client/server models → message passing primitives
- › Shared memory for reducing the overhead

Some special OS supports only one of the two

- › RTOS supports only shared memory





# Models of concurrency



# Message passing

---

Message passing systems are based on the basic concept of message

Two basic operations

`send (destination, message)`

- ✓ send can be synchronous or asynchronous (*fire-and-forget*)

`receive (source, &message)`

- ✓ receive can be symmetric or asymmetric



# Producer/Consumer with MP

---

- ✓ The producer executes `send (consumer, data)`
- ✓ the consumer executes `receive (producer, data)`
- ✓ no need for a special communication structure (already contained in the send/receive semantic)





# Resources and message passing

---

There are no shared resources in the message passing model

- ✓ all the resources are allocated statically, accessed in a dedicated way

Each resource is handled by a manager process that is the only one that has right to access to a resource

- ✓ The consistency of a data structure is guaranteed by the **manager** process
- ✓ There is no more competition, only cooperation!!!



# Synchronous communication

Synchronous send/receive



✓ no buffers!

producer:

```
s_send (consumer, d);
```

consumer:

```
s_receive (producer, &d);
```





# Async send/ sync receive

Asynchronous send / synchronous receive


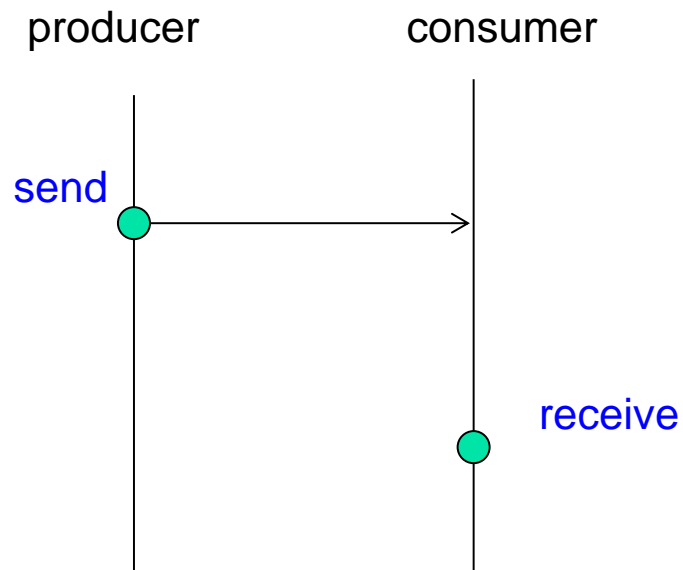
- ✓ there is probably a send buffer somewhere

producer:

```
a_send (consumer, d);
```

consumer:

```
s_receive (producer, &d);
```





# (A)symmetric receive

---

**Symmetric** receive                      `receive (source, &data);`

- ✓ the programmer wants a message from a given producer

**Asymmetric** receive                      `source = receive (&data);`

- ✓ often, we do not know who is the sender

E.g., a web server is asymmetric


- ✓ the programmer cannot know in advance the address of the browser that will request the service
- ✓ many browsers can ask for the same service



# Remote procedure call

From low-level (MP) to more «programmer friendly» (RPC) mechanism

- ✓ Increase expressiveness, higher level of abstraction
- ✓ In a client-server system, a client wants to request an action to a server
- ✓ Typically done using a remote procedure call (RPC)







# Message passing systems




In message passing

- ✓ each resource needs one threads manager (often called **daemon thread**)
- ✓ the threads manager is responsible for giving access to the resource

Example: mutual exclusion with message passing primitives

- ✓ one thread will ensure mutual exclusion
- ✓ Every thread that wants to access the resource must
  - Send a message to the manager thread
  - Access the critical section
  - Send a message to signal the leaving of the critical section

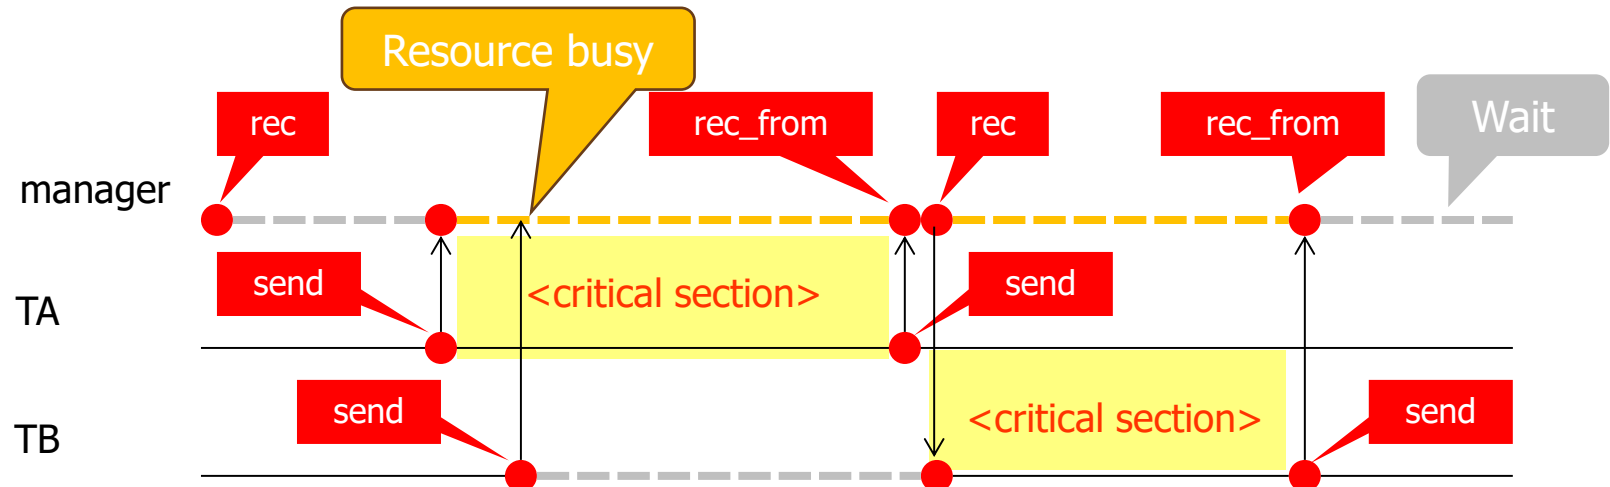




# Sync send / sync receive

```
void * manager(void *) {  
    thread_t source;  
    int d;  
    while (true) {  
        source = s_receive (&d);  
        s_receive_from (source, &d);  
    }  
}
```

```
void * thread(void *) {  
    int d;  
    while (true) {  
        s_send(manager, d);  
  
        <critical section>  
  
        s_send(manager, d);  
    }  
}
```





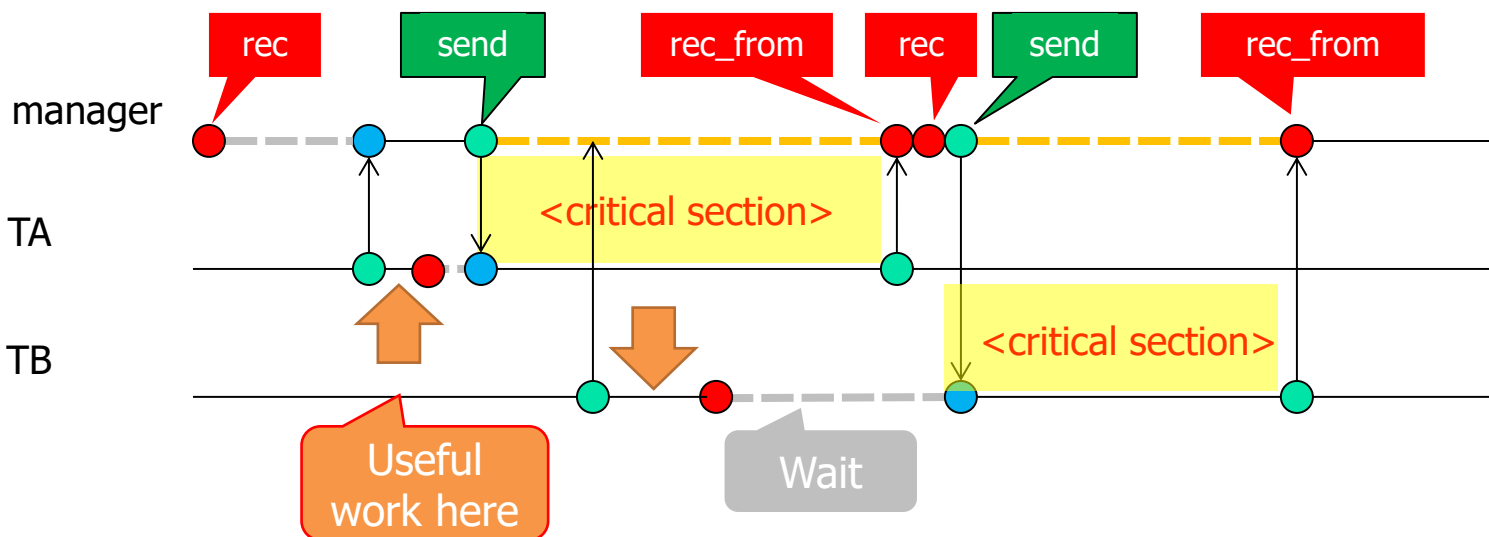
# With async send and sync receive

```
void * manager(void *) {  
    thread_t source;  
    int d;  
    while (true) {  
        source = s_receive (&d);  
        a_send (source, &d);  
        s_receive_from (source, &d);  
    }  
}
```

Blocking

```
void * thread(void *) {  
    int d;  
    while (true) {  
        a_send (manager, d);  
        s_receive_from (manager, d);  
  
        <critical section>  
  
        a_send (manager, d);  
    }  
}
```

Non blocking



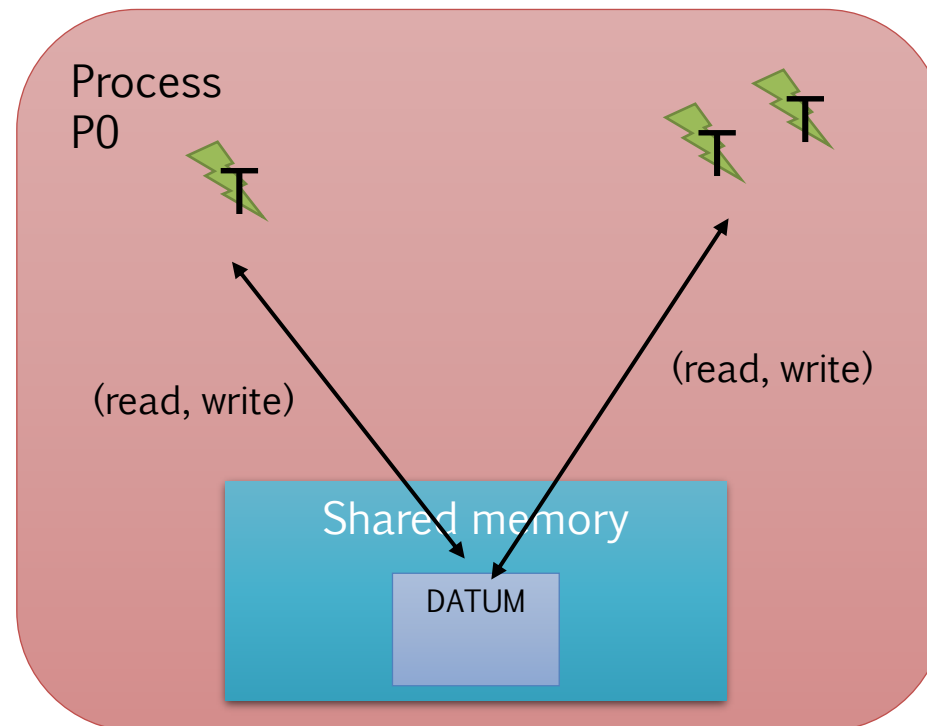


Shared memory model



# Shared memory abstraction

- ✓ The first one being supported in old OS
- ✓ The simplest one and the closest to the machine
- ✓ All threads can access the same memory locations

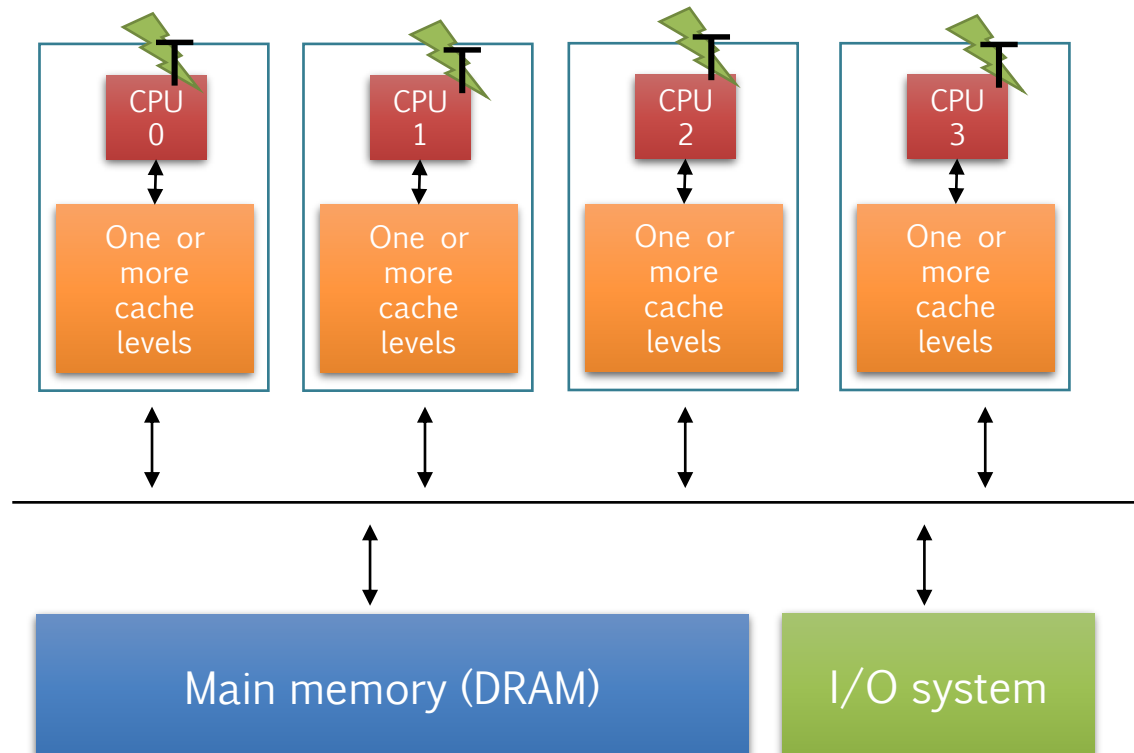




# Analogy with hardware

An **abstract** model that presents a good analogy is the following:

- › Many HW CPU, each one running one activity (thread)
- › One shared memory





# Resource allocation

---

Allocation of resource can be

- ✓ **Static:** once the resource is granted, it is never revoked
- ✓ **Dynamic:** resource can be granted and revoked dynamically
  - Manager

Access to a resource can be

- ✓ **Dedicated:** only one activity at a time may request access to the resource
- ✓ **Shared:** many activities may access the resource at the same time
  - Mutual exclusion

	Dedicated	Shared
Static	Compile Time	Manager
Dynamic	Manager	Manager



# Mutual exclusion: a (big) problem

---

We do not know in advance the relative speed of the processes

- ✓ Hence, we do not know the order of execution of the hardware instructions

Example:

- ✓ Incrementing a variable  $x$  is NOT an atomic operation





# Atomicity

---

A hardware instruction is atomic if it cannot be “interleaved” with other instructions

- ✓ Atomic operations are always sequentialized

Atomic operations cannot be interrupted

- ✓ They are thread safe operations
  - For example, transferring one word from memory to register or viceversa

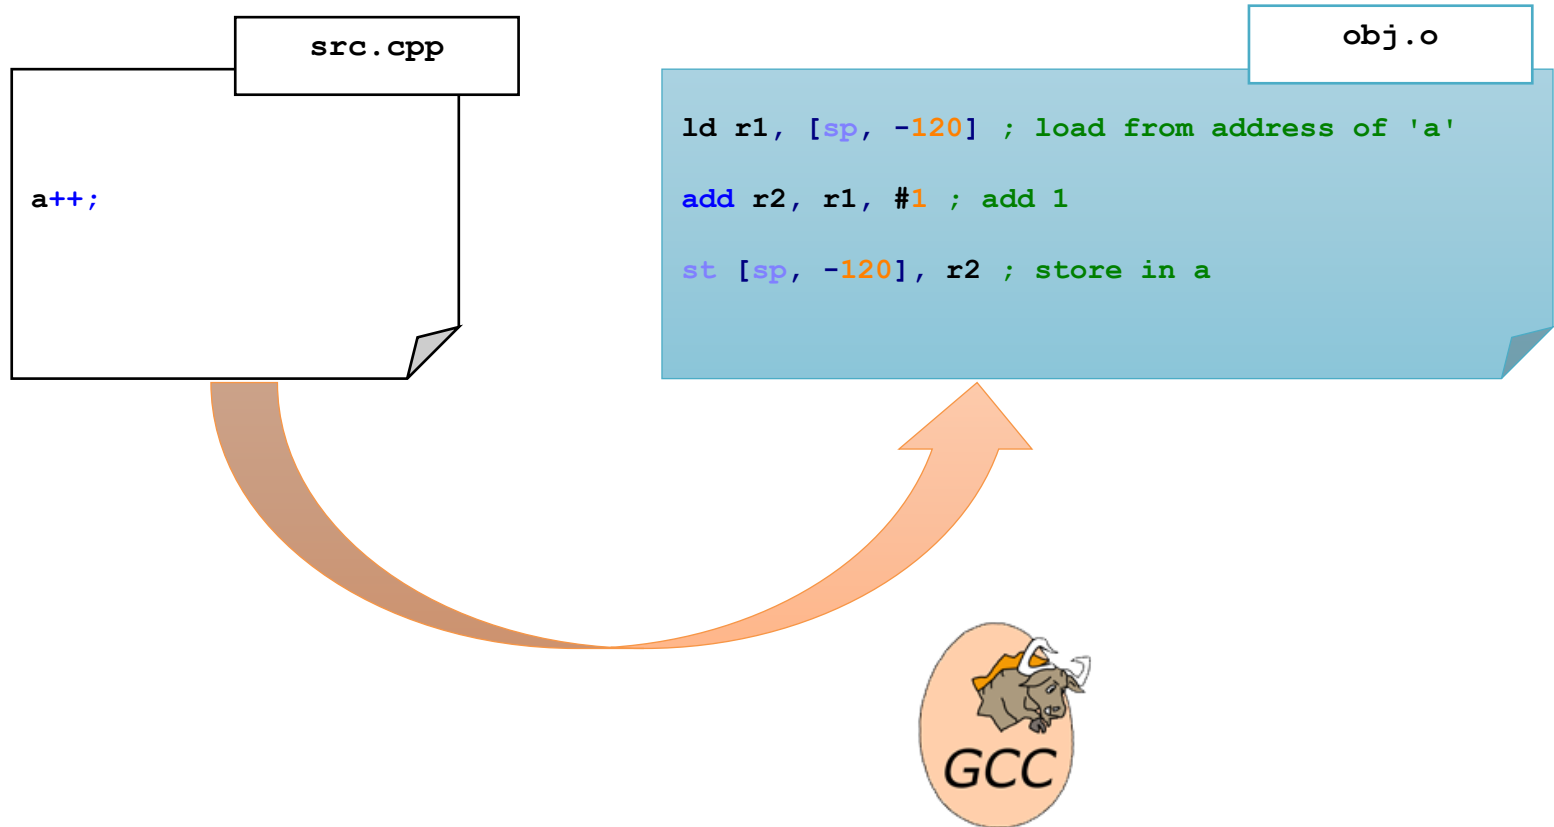
Non atomic operations can be interrupted

- ✓ They are not “safe” operations
- ✓ Non-elementary operations are not atomic



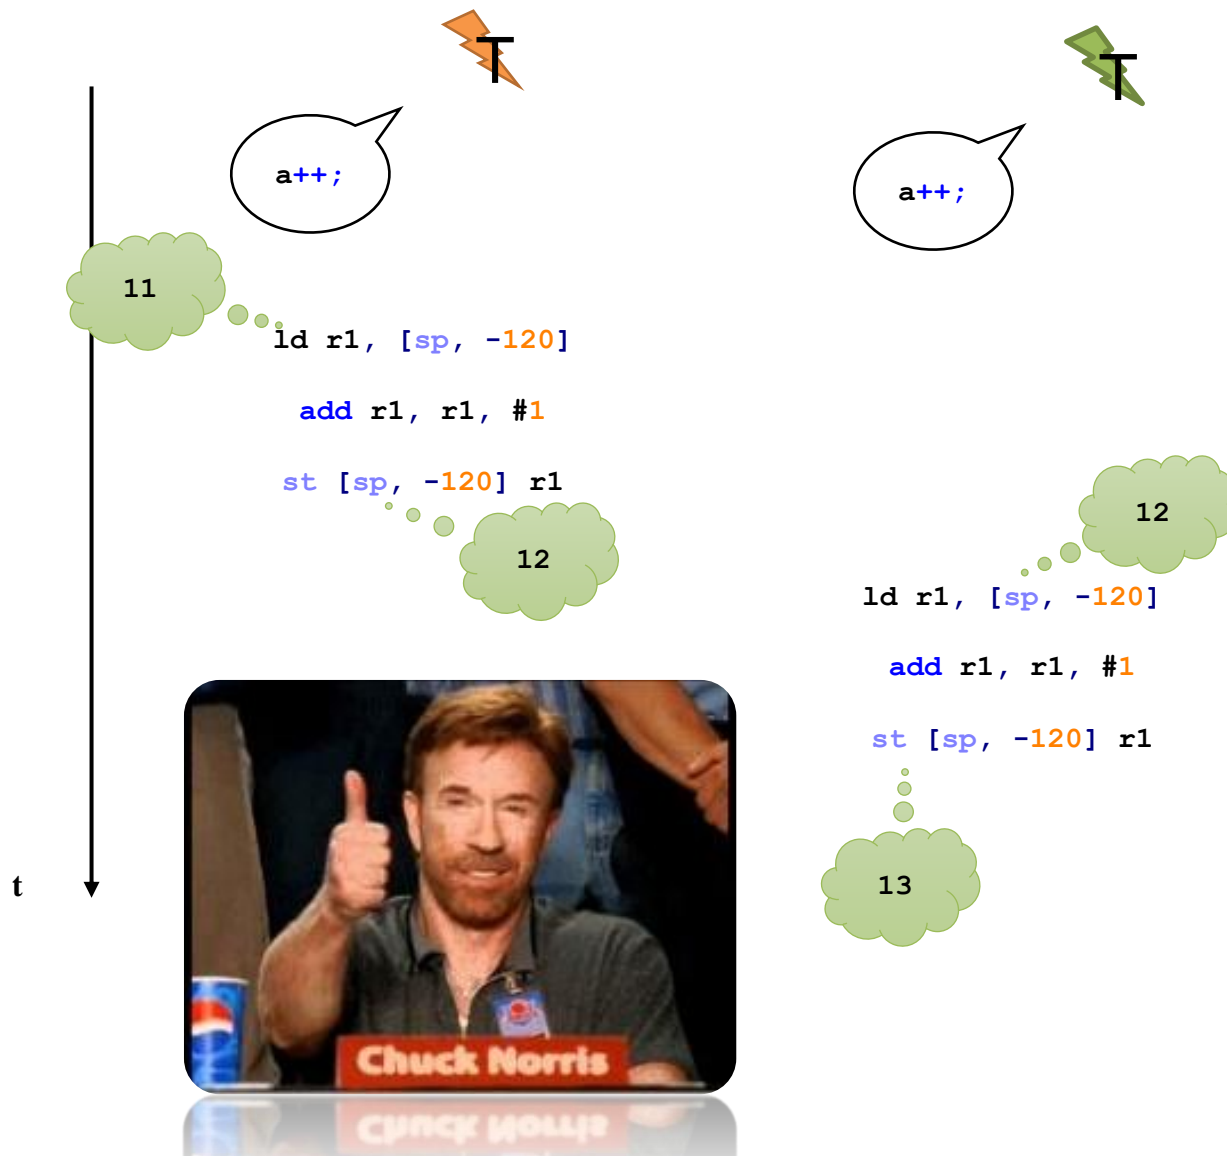
# The big problem

- › a++ is not an unique instruction



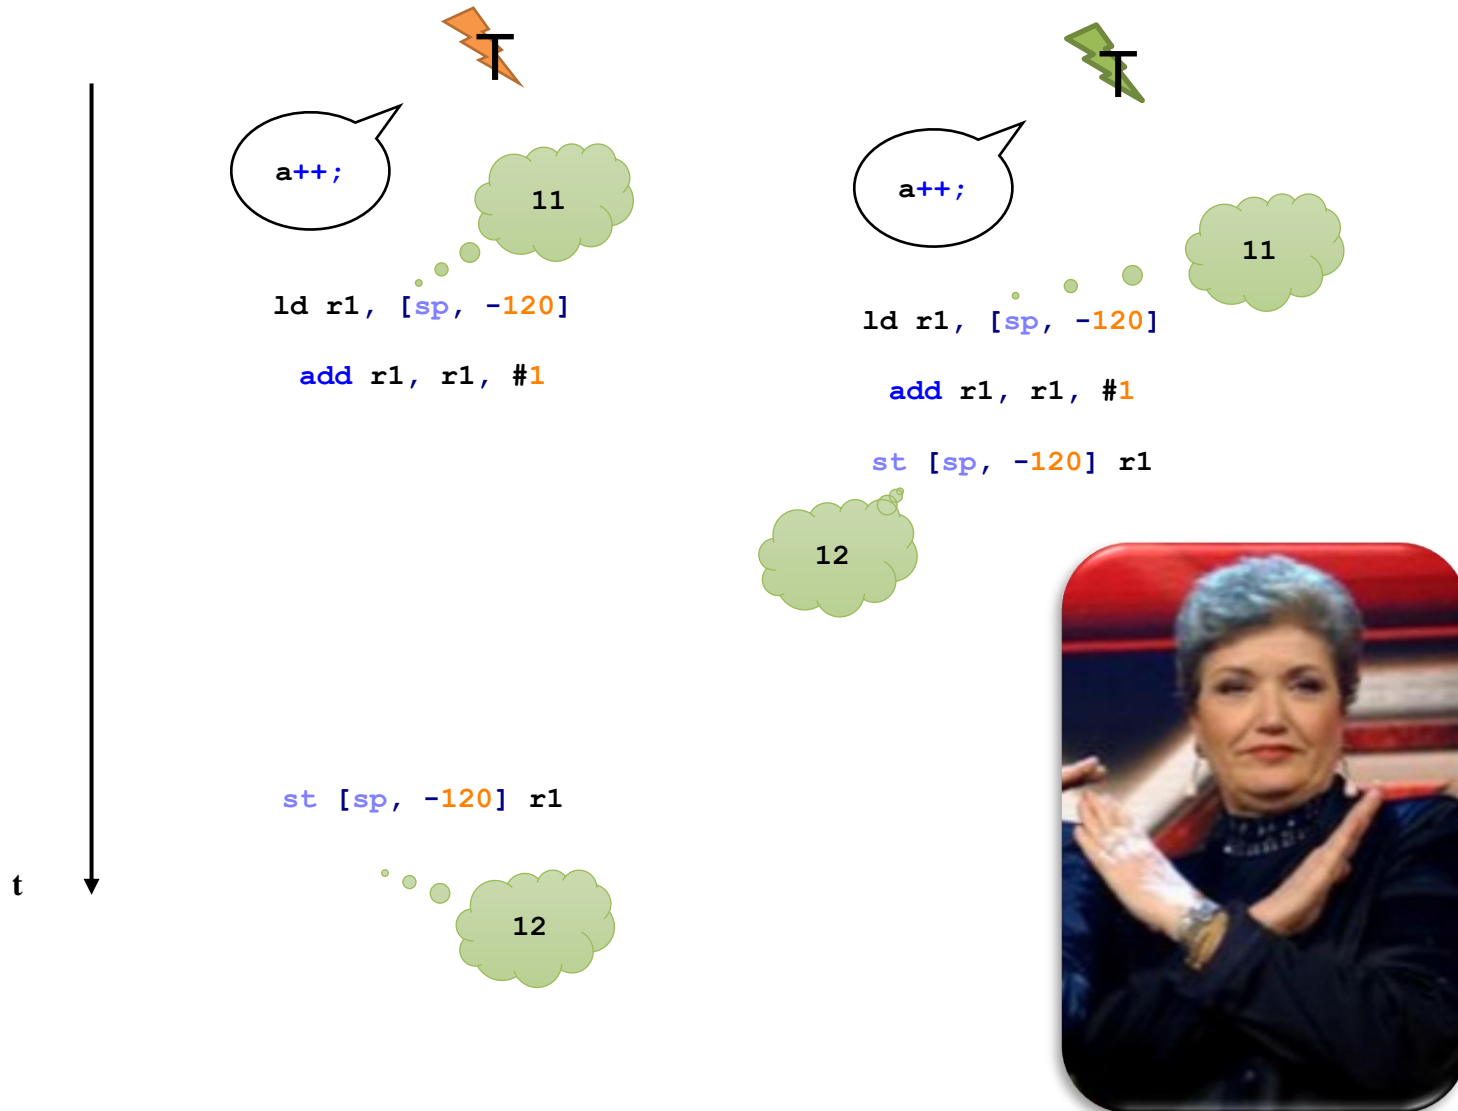


# When things go well





# When things go less well





# Critical sections

---

## Definitions

- ✓ The shared object (e.g.,  $x$ ) where the conflict may happen is a “**resource**”
- ✓ The parts of the code where the problem may happen are called “**critical sections**”

A critical section is a sequence of operations that cannot be interleaved with other operations on the same resource

Multiple critical sections on the same resource must execute in **MUTUAL EXCLUSION**

- ✓ atomic operation (hardware extension within the memory banks)
- ✓ semaphores
- ✓ mutexes
- ✓ *locks*



# General mechanism: semaphores

---

Proposed by Dijkstra

A semaphore is an abstract entity that consists of

- ✓ A counter
- ✓ A blocking queue (of threads)

Can perform two atomic operations

- ✓ Blocking **Wait** for a given condition
- ✓ **Signal** that the condition becomes true (aka **Post**)

We can also use them to implement mutual exclusion (we'll see this)



# Wait and Signal

---

A **Wait** operation has the following behavior

- ✓ If `counter == 0`, the requiring thread is blocked
  - It is removed from the ready queue
  - It is inserted in the blocked queue
- ✓ If `counter > 0`, then `counter--`;

A **Signal (aka: Post)** operation has the following behavior

- ✓ If `counter == 0` and there is some blocked thread, unblock it
  - The thread is removed from the blocked queue
  - It is inserted in the ready queue
- ✓ Otherwise, increment counter



# Semaphores (pseudo-code)

```
void sem_init (sem_t *s, int n) {
    s->count=n;
    // ...
}

void sem_wait(sem_t *s) {
    if (s->count == 0)
        <block the thread>
    else
        s->count--;
}

void sem_post(sem_t *s) {
    if (<there are blocked threads>)
        <unlock a thread>
    else
        s->count++;
}
```

- › We assume that `s->count` (and its operation `++`) are atomic





# Signal semantics

---



What happens when a thread blocks on a semaphore?

- ✓ In general, it is inserted in a BLOCKED queue

Extraction from the blocking queue can follow different semantics:

- ✓ Strong semaphore
  - The threads are removed in well-specified order
  - For example, FIFO order, priority based ordering, ...
- ✓ Signal and suspend
  - After the new thread has been unblocked, a thread switch happens
- ✓ Signal and continue
  - After the new thread has been unblocked, the thread that executed the signal continues to execute

Concurrent programs should not rely too much on the semaphore semantic



# Mutual exclusion with semaphores

How to use a semaphore for critical sections?

- ✓ Define a semaphore **initialized to 1**
- ✓ Before entering the critical section, perform a **wait**
- ✓ After leaving the critical section, perform a **signal/post**

```
sem_t s;  
  
// ...  
sem_init (&s, 1);
```

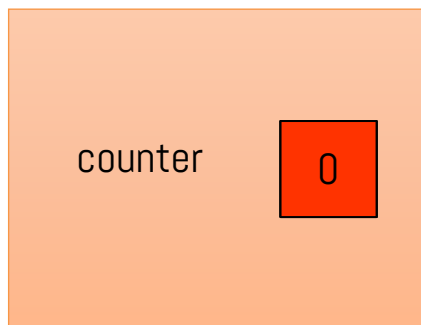
```
void *threadA(void *arg) {  
    sem_wait (&s);  
  
    <critical section>  
  
    sem_post (&s);  
}
```

```
void *threadB(void *arg) {  
    sem_wait (&s);  
  
    <critical section>  
  
    sem_post (&s);  
}
```



# Mutual exclusion with semaphores

semaphore



(TA) sem\_wait();  
(TA) <critical section (1)>



(TA) <critical section (2)>  
(TA) sem\_post()

sem\_wait() (TB)

<critical section> (TB)  
sem\_post() (TB)

$t$



Let's see  
this in  
action





# Synchronisation


---

Mutual exclusion is not the only problem: we need a way of synchronise two or more threads

✓ Example: producer/consumer

We have two threads,

- ✓ One produces some integers and sends them to another thread (**PRODUCER**)
- ✓ Another one takes the integer and elaborates it (**CONSUMER**)





# Synchronization with semaphores

Define a semaphore initialized to 0 (**blocked**)

- ✓ At the synchronization point, consumer performs a **Wait**
- ✓ At the synchronization point, producer performs a **Signal/Post**
- ✓ In the example, *threadA* blocks until *threadB* wakes it up

```
sem_t s;  
  
// ...  
sem_init (&s, 0);
```

```
void *threadA(void *arg) {  
    sem_wait (&s);  
    process (a);  
}
```

<<unlock>>

```
void *threadB(void *arg) {  
    a = 11;  
    sem_post (&s);  
}
```



# Producer/consumer: how to do it naively

Share a queue of data/objects/anything you might need to produce&consume

- ✓ If the queue is full, the producer actively waits
- ✓ If the queue is empty, the consumer actively waits
- ✓ Aka: **busy-waiting**

Very inefficient!

```
struct CircularArray_t queue;
```

```
void *producer(void *) {  
    bool res;  
    int data;  
    while(1) {  
        <obtain data>  
        while (!insert_CA(&queue,  
                           data))  
            ;  
    }  
}
```

```
void *consumer(void *) {  
    bool res;  
    int data;  
    while(1) {  
        while (!extract_CA(&queue,  
                           &data))  
            ;  
        <use data>  
    }  
}
```



# Naive, polling-based producer/consumer

---

Consider a producer/consumer system

Producer(s) execute `insert_CA()`

- ✓ We want the producers to be blocked when the queue is full
- ✓ The producers will be unblocked when there is some space again

Consumer(s) execute `extract_CA()`

- ✓ We want the consumers to be blocked when the queue is empty
- ✓ The consumers will be unblocked when there is some space again
- ✓ First attempt: one producer and one consumer only



# One producer, one consumer

```
struct CircularArray_t {
    int array[10];
    int head, tail;
    sem_t empty, full;
}

void init_CA (struct CircularArray_t *c) {
    c->head=0; c->tail=0;
    sem_init(&c->empty, 0); sem_init(&c->full, 10);
}

void insert_CA (struct CircularArray_t *c, int elem) {
    sem_wait(&c->full);
    c->array[c->head] = elem;
    c->head = (c->head + 1) % 10;
    sem_post(&c->empty);
}

void extract_CA (struct CircularArray_t *c, int &elem) {
    sem_wait(&c->empty);
    elem = c->array[c->tail];
    c->tail = (c->tail + 1) % 10;
    sem_post(&c->full);
}
```

*Block if queue is full*

*Release those who're waiting for extraction*

*Block if queue is empty*

*Release those who're waiting for insertion*





# Multiple producers/consumers

---

Combine mutual exclusion and synchronization

- ✓ Semaphore to implement synchronization
- ✓ Semaphore to protect the data structure
  - ✓ Make the primitives (`insert_CA`, `extract_CA`) THREAD-SAFE



# Producers/consumers: does this work?

```
struct CircularArray_t {
    int array[10];
    int head, tail;
    sem_t full, empty;
    sem_t mutex;
}

void init_CA(struct CircularArray_t *c) {
    c->head=0; c->tail=0;
    sem_init (&c->empty, 0); sem_init (&c->full, 10); sem_init (&c->mutex, 1);
}
```

```
void insert_CA (struct CircularArray_t *c,
               int elem) {
    sem_wait (&c->mutex);
    sem_wait (&c->full);
    c->array[c->head]=elem;
    c->head = (c->head+1)%10;
    sem_post (&c->empty);
    sem_post (&c->mutex);
}
```

*Enter  
critical  
section*

*Exit  
critical  
section*

```
void extract_CA (struct CircularArray_t *c,
                int *elem) {
    sem_wait (&c->mutex);
    sem_wait (&c->empty);
    *elem = c->array[c->tail];
    c->tail = (c->tail+1)%10;
    sem_post (&c->full);
    sem_post (&c->mutex);
}
```

...of course NOT!

- › Why? (red => protects critical section/mutual exclusion;  
other colors => synchronization - producer/consumer)



# The deadlock explained

---

- ✓ A thread executes `sem_wait(&c->mutex)` and then blocks on a synchronisation semaphore
- ✓ To be unblocked another thread must enter a critical section guarded by the same mutex semaphore!
- ✓ So, the first thread cannot be unblocked and free the mutex!

The situation cannot be solved, and the two threads will never proceed

As a rule, **never insert a blocking synchronization** inside a critical section!!!



# Producers/consumers: correct solution

```
struct CircularArray_t {
    int array[10];
    int head, tail;
    sem_t full, empty;
    sem_t mutex;
}

void init_CA(struct CircularArray_t *c) {
    c->head=0; c->tail=0;
    sem_init (&c->empty, 0); sem_init (&c->full, 10); sem_init (&c->mutex, 1);
}
```

```
void insert_CA (struct CircularArray_t *c,
               int elem) {
    sem_wait (&c->full);
    sem_wait (&c->mutex);
    c->array[c->head]=elem;
    c->head = (c->head+1)%10;
    sem_post (&c->mutex);
    sem_post (&c->empty);
}
```

```
void extract_CA (struct CircularArray_t *c,
                int *elem) {
    sem_wait (&c->empty);
    sem_wait (&c->mutex);
    elem = c->array[c->tail];
    c->tail = (c->tail+1)%10;
    sem_post (&c->mutex);
    sem_post (&c->full);
}
```

...we were mixing producer/consumer and synchronizastion

› NEVER NEST MUTEXES



## Course website

- › [http://hipert.unimore.it/people/paolob/pub/Industrial\\_Informatics/index.html](http://hipert.unimore.it/people/paolob/pub/Industrial_Informatics/index.html)

## My contacts

- › [paolo.burgio@unimore.it](mailto:paolo.burgio@unimore.it)
- › <http://hipert.mat.unimore.it/people/paolob/>

## Resources

- › Giorgio Buttazzo, "Hard Real-Time Computing Systems : Predictable Scheduling Algorithms and Applications". 3<sup>rd</sup> Edition. 2011. Springer
- › "Real-Time Embedded Systems" course by Prof. Bertogna @UNIMORE
- › A "small blog"
  - <http://www.google.com>