

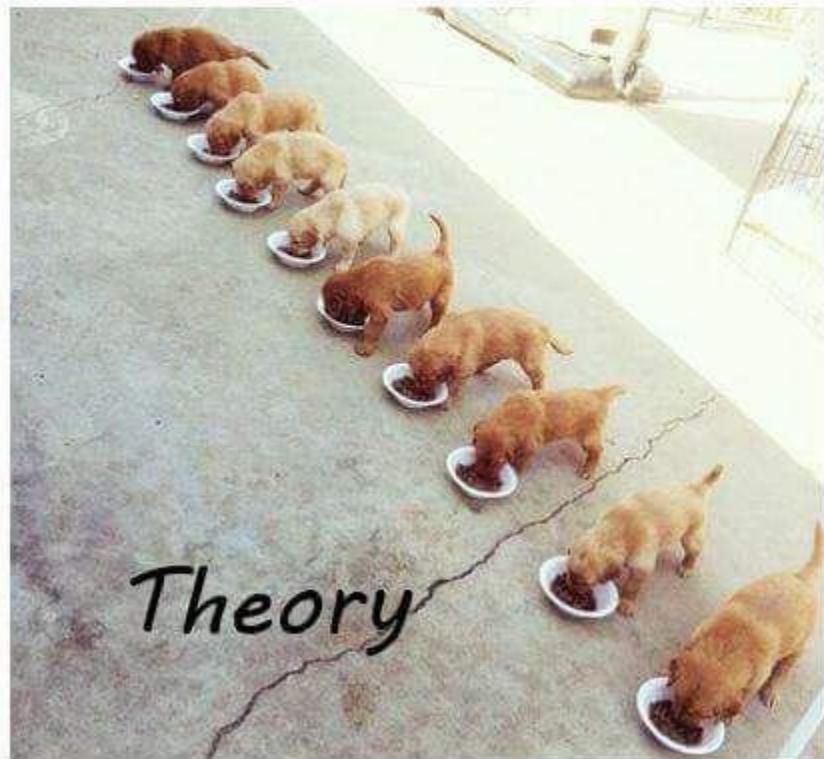
# General Purpose GPUs

## GP-GPUs



Paolo Burgio  
[paolo.burgio@unimore.it](mailto:paolo.burgio@unimore.it)

# Multithreaded programming



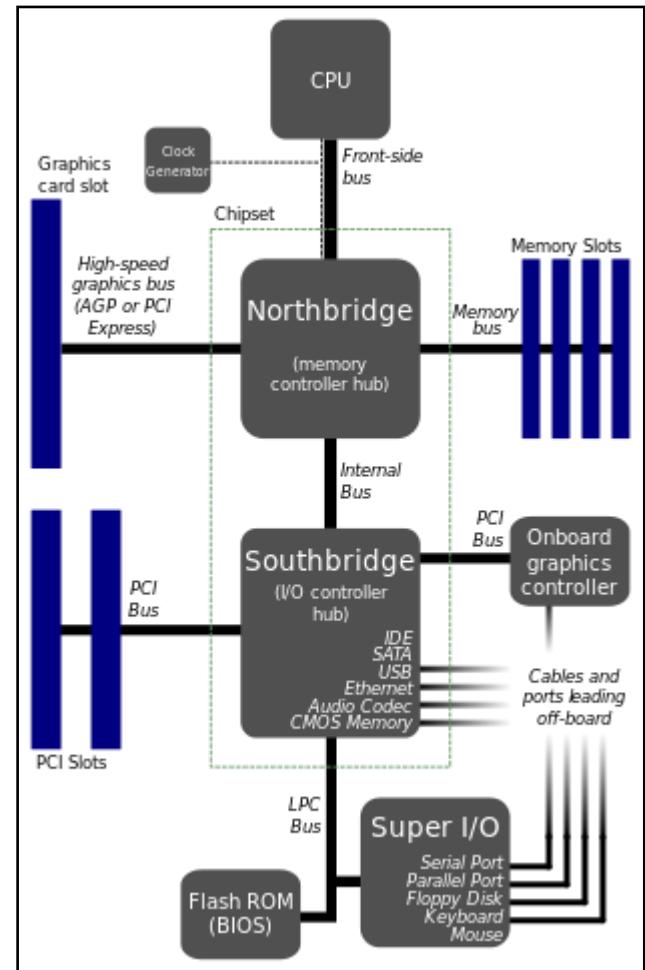
Theory



Actual

# Graphics Processing Units

- › (Co-)processor devoted to graphics
  - Built as "monolithical" chip
  - Integrated as co-processor
  - Recently, SoCs
- › Main providers
  - NVIDIA
  - ATI
  - AMD
  - Intel...
- › We will focus on NVIDIA
  - Widely adopted
  - Adopted by us



# A bit of history...

- › 70s: first "known" graphic card on a board package
- › Early 90s: 3D graphics popular in **games**
- › 1992: **OpenGL**
- › 1999: NVIDIA GeForce 256 "World's first GPU"
- › 2001: NVIDIA GeForce 3, w/programmable shaders (First **GP-GPU**)
- › 2008: NVIDIA GeForce 8800 GTX w/**CUDA** capabilities - Tesla arch.
- › 2009: **OpenCL 1.0** inside MAC OS X Snow Leopard
- › 2010: NVIDIA GeForce 400 Series - Fermi arch.
- › 2010-1: OpenCL 1.1, 1.2
- › 2012: NVIDIA GeForce 600 Series - Kepler arch.
- › 2013: OpenCL 2.0
- › 2014: NVIDIA GeForce 745 OEM - Maxwell arch.
- › **2015 Q4: NVIDIA and HiPeRT Lab start cooperation ;)**
- › 2017 Q1: NVIDIA Drive Px2 for Self-Driving Cars
- › 2019 Q1: NVIDIA Pegasus for Self-Driving Cars





# ...a bit of confusion!

---

---

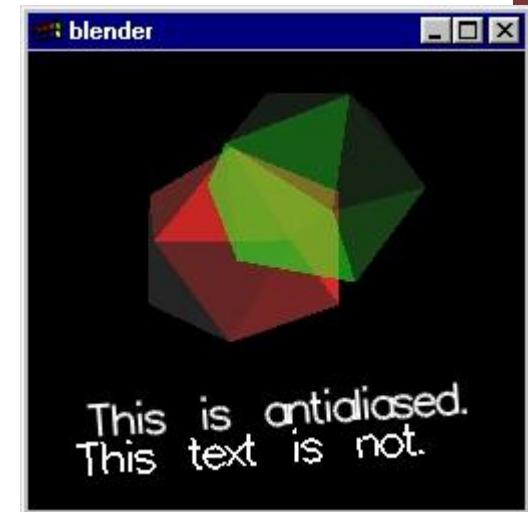
- › Many architectures
  - Tesla, Fermi, Maxwell, Pascal, Volta..
- › Many programming librar... langag... frameworks
  - OpenGL
  - CUDA
  - OpenCL
  - ...
- › Many application domains!
  - Graphics
  - GP-GPUs?
  - Automotive!??!?!??!
- › Let's start from scratch...



# GPU for graphics - OpenGL

- › Use GPUs for rendering of graphics

- A library of functions and datatypes
- Use directly in the code
- High-level operations on lights, shapes, shaders...



- › Tailored for the specific domain and **programmer skills**

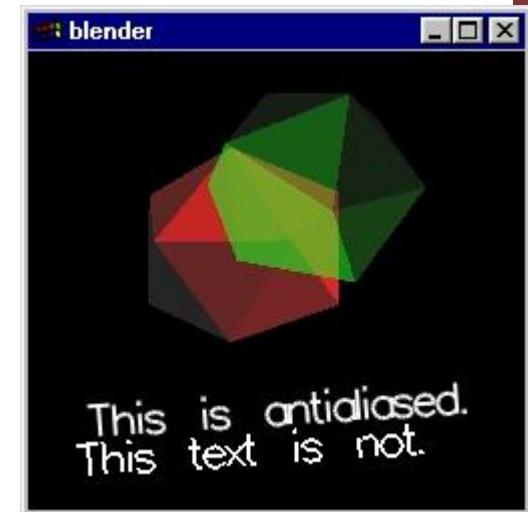
- Hides away the complexity of the machine
- Takes care of "low" level optimizations/operations



# GPU for graphics - OpenGL

- › Use GPUs for rendering of graphics

- A library of functions and datatypes
- Use directly in the code
- High-level operations on lights, shapes, shaders...



- › Tailored for the specific domain and **programmer skills**

- Hides away the complexity of the machine
- Takes care of "low" level optimizations/operations

```
int main(int argc, char **argv) {  
    glutInit(&argc, argv);  
    glutInitDisplayMode(GLUT_DOUBLE | GLUT_RGB | GLUT_DEPTH);  
    glutCreateWindow("blender");  
    glutDisplayFunc(display);  
    glutVisibilityFunc(visible);  
  
    glNewList(1, GL_COMPILE); /* create ico display list */  
    glutSolidIcosahedron();  
    glEndList();  
  
    glEnable(GL_LIGHTING);  
    glEnable(GL_LIGHT0);  
    glLightfv(GL_LIGHT0, GL_AMBIENT, light0_ambient);  
    glLightfv(GL_LIGHT0, GL_DIFFUSE, light0_diffuse);  
    glLightfv(GL_LIGHT1, GL_DIFFUSE, light1_diffuse);  
    glLightfv(GL_LIGHT1, GL_POSITION, light1_position);  
    glLightfv(GL_LIGHT2, GL_DIFFUSE, light2_diffuse);  
    glLightfv(GL_LIGHT2, GL_POSITION, light2_position);  
    glEnable(GL_DEPTH_TEST);  
    glEnable(GL_CULL_FACE);  
    glEnable(GL_BLEND);  
    glBlendFunc(GL_SRC_ALPHA, GL_ONE_MINUS_SRC_ALPHA);  
    glEnable(GL_LINE_SMOOTH);  
  
    glLineWidth(2.0);  
    glMatrixMode(GL_PROJECTION);  
    gluPerspective(/* field of view in degree */ 40.0,  
                  /* aspect ratio */ 1.0,  
                  /* Z near */ 1.0,  
                  /* Z far */ 10.0);  
    glMatrixMode(GL_MODELVIEW);  
    gluLookAt(0.0, 0.0, 5.0, /* eye is at (0,0,5) */  
              0.0, 0.0, 0.0, /* center is at (0,0,0) */  
              0.0, 1.0, 0.); /* up is in positive Y direction */  
    glTranslatef(0.0, 0.6, -1.0);  
  
    glutMainLoop();  
    return 0; /* ANSI C requires main to return int. */  
}
```

# OpenGL

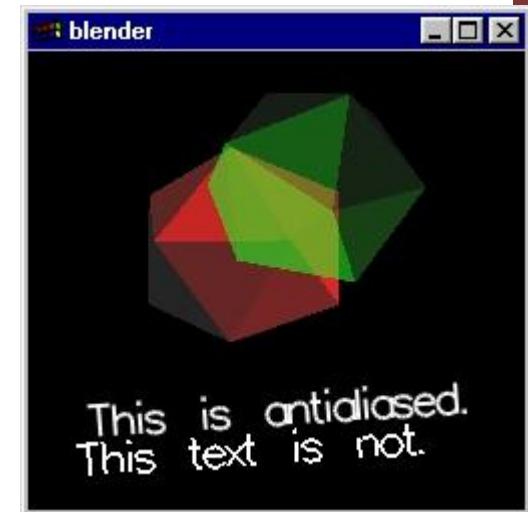


ammer skills

# GPU for graphics - OpenGL

## › Use GPUs for rendering of graphics

- A library of functions and datatypes
- Use directly in the code
- High-level operations on lights, shapes, shaders...



## › Tailored for the specific domain and ~~programmer skills~~

- Hides away the complexity of the machine
- Takes care of "low" level optimizations/operations

```
GLfloat light0_ambient[] = {0.2, 0.2, 0.2, 1.0};  
GLfloat light0_diffuse[] = {0.0, 0.0, 0.0, 1.0};  
GLfloat light1_diffuse[] = {1.0, 0.0, 0.0, 1.0};  
GLfloat light1_position[] = {1.0, 1.0, 1.0, 0.0};  
GLfloat light2_diffuse[] = {0.0, 1.0, 0.0, 1.0};  
GLfloat light2_position[] = {-1.0, -1.0, 1.0, 0.0};
```



# General Purpose - GPUs

---

---

- › We have a machine with thousand of cores
  - why should we use it only for graphics?
- › Use it for General Purpose Computing!
  - GP-GPU
  - ~yr 2000

*NdA:* Computing modes

- General Purpose Computing
- High-Performance Computing
- Embedded Computing
- Real-Time Computing
- ...

# General Purpose - GPUs

---

---

- › We have a machine with thousand of cores
  - why should we use it only for graphics?
- › Use it for General Purpose Computing!
  - GP-GPU
  - ~yr 2000



*NdA:* Computing modes

- General Purpose Computing
- High-Performance Computing
- Embedded Computing
- Real-Time Computing
- ...

# General Purpose - GPUs

- > We have a  
– why should we
- > Use it for
  - GP-GPU
  - ~yr 2000



*NdA:* Computing modes

- General Purpose Computing
- High-Performance Computing
- Embedded Computing
- Real-Time Computing
- ...

# General Purpose - GPUs

- › We have a machine with thousand of cores
  - why should we use it only for graphics?
- › Use it for General Purpose Computing!
  - GP-GPU
  - ~yr 2000



## NdA: Computing modes

- General Purpose Computing
- High-Performance Computing
- Embedded Computing
- Real-Time Computing
- ...

# General Purpose - GPUs

- › We have a machine with thousand of cores
  - why should we use it only for graphics?

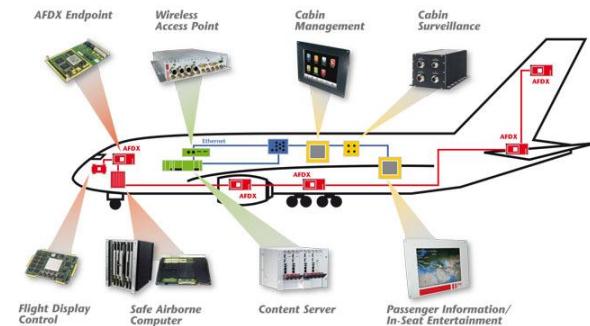


- › Use it for General
  - GP-GPU
  - ~yr 2000



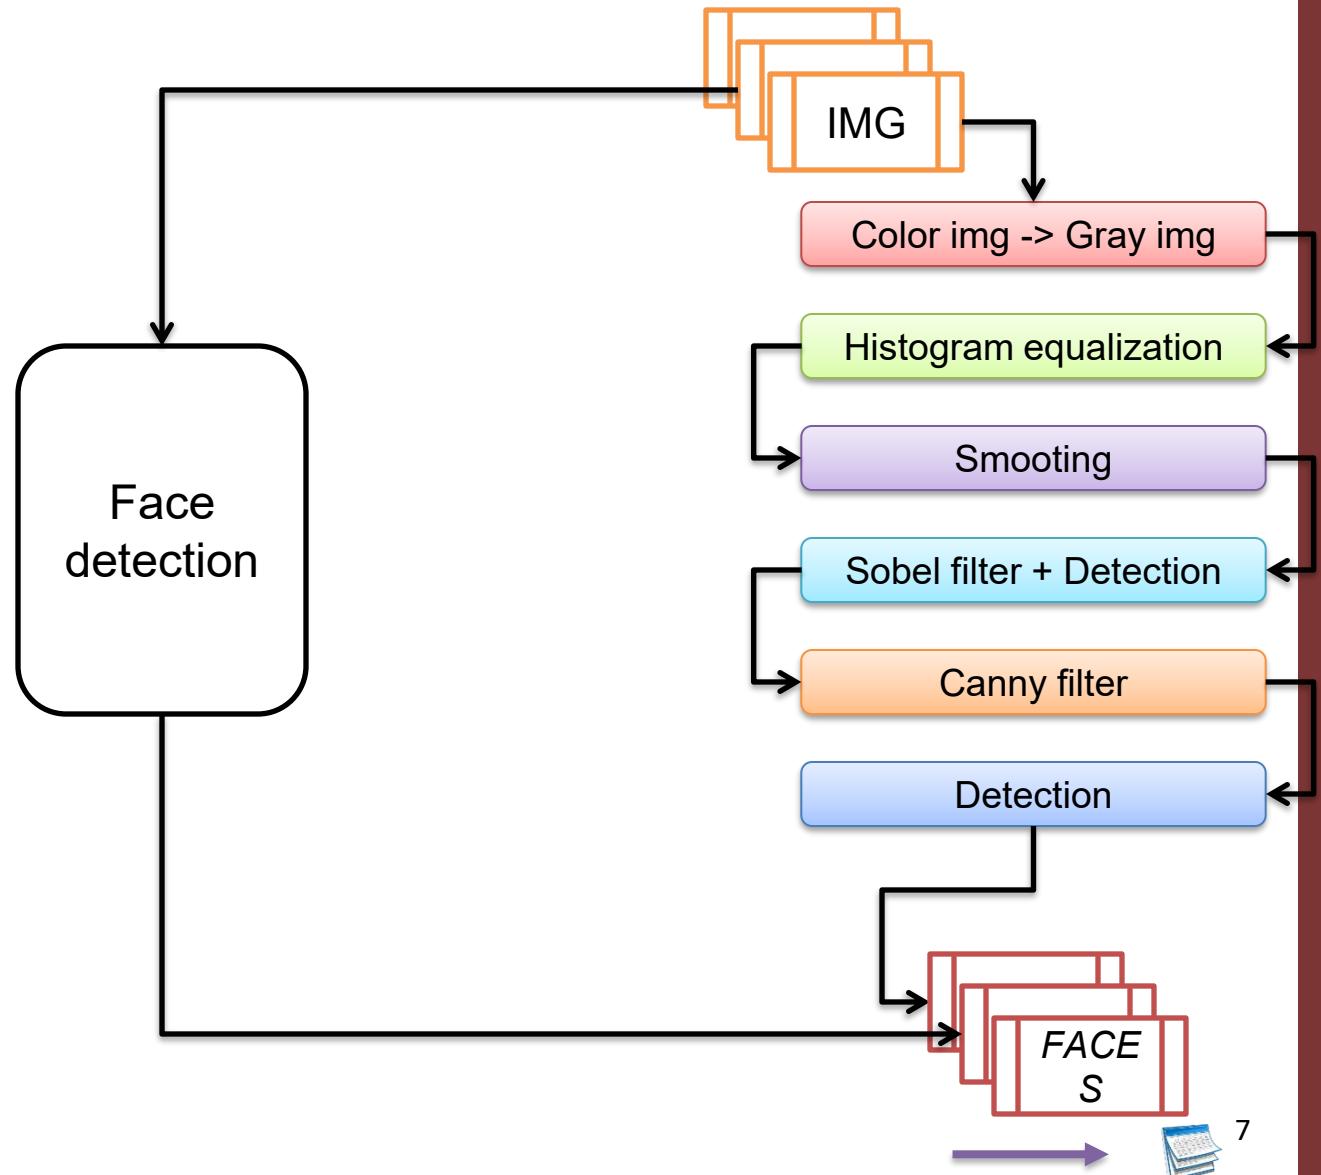
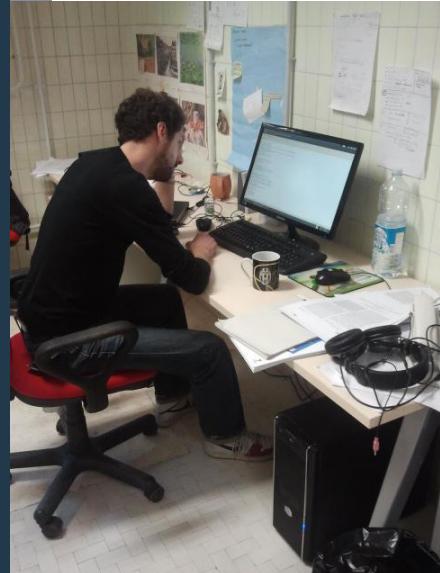
## NdA: Computing modes

- General Purpose Computing
- High-Performance Computing
- Embedded Computing
- Real-Time Computing
- ...





# Under the hood: face detection





# Under the hood: face detection



Face  
detection



IMG

Color img -> Gray img

Histogram equalization

Smoothing

Sobel filter + Detection

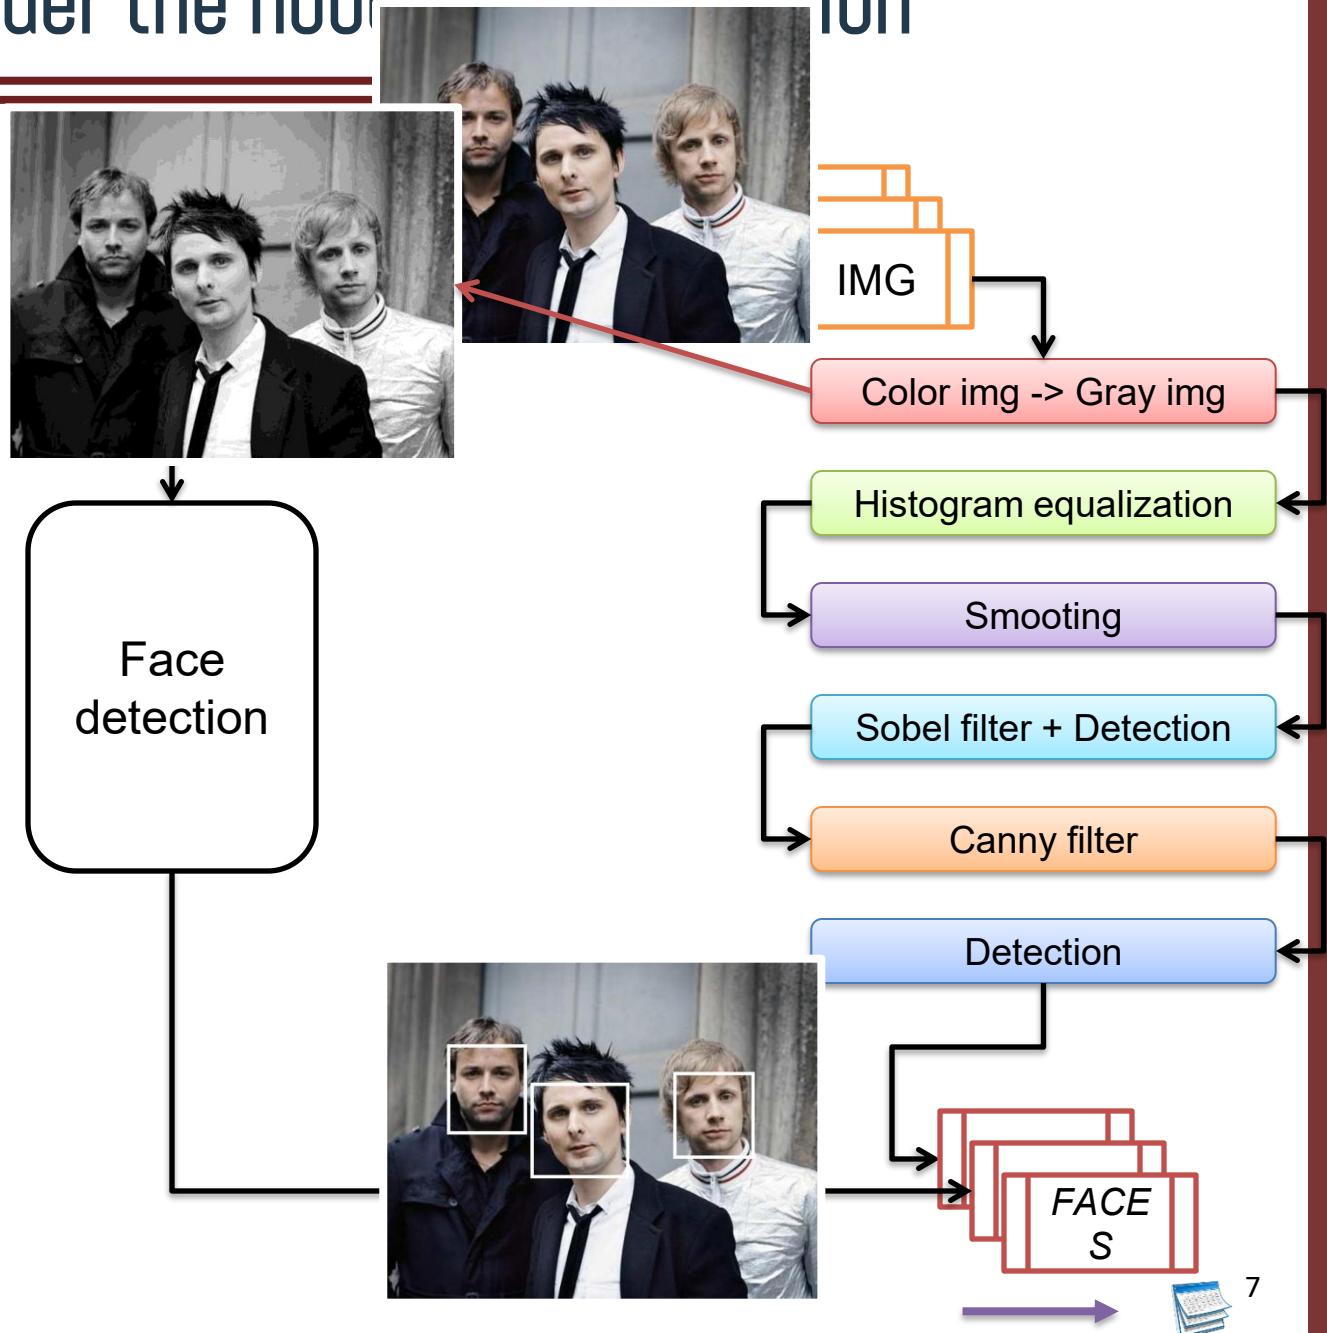
Canny filter

Detection

FACE  
S

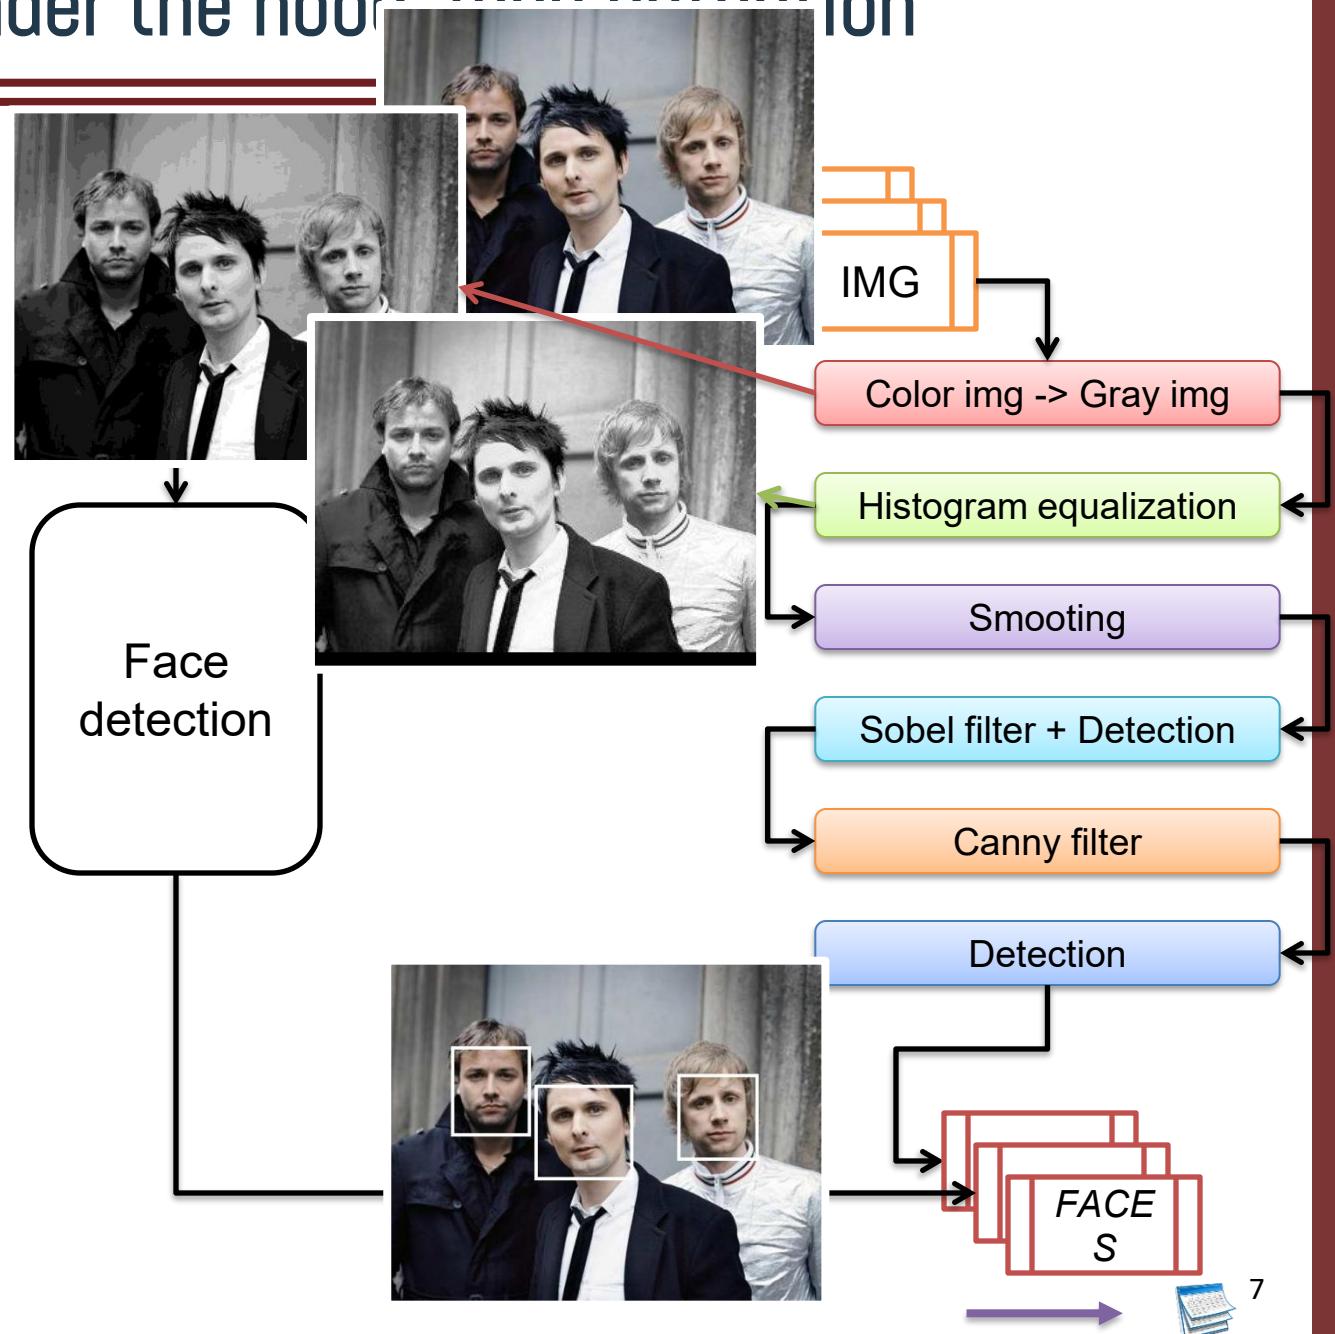


# Under the hood: face detection



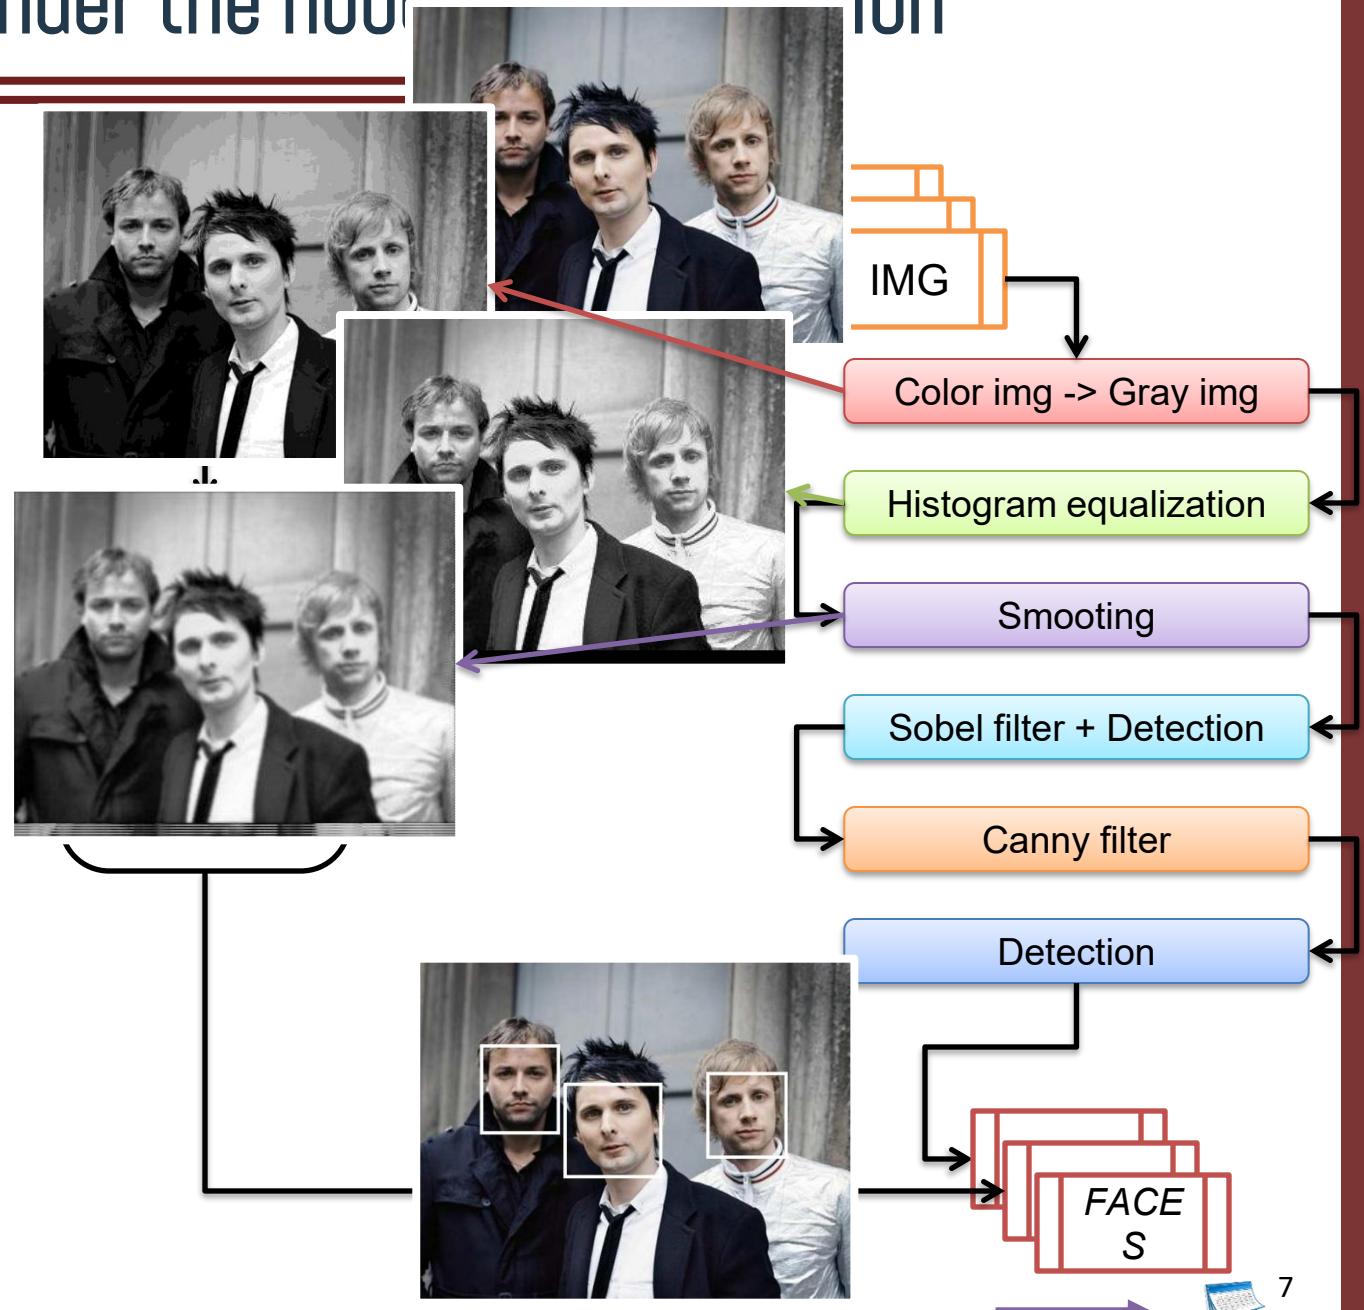


# Under the hood: face detection

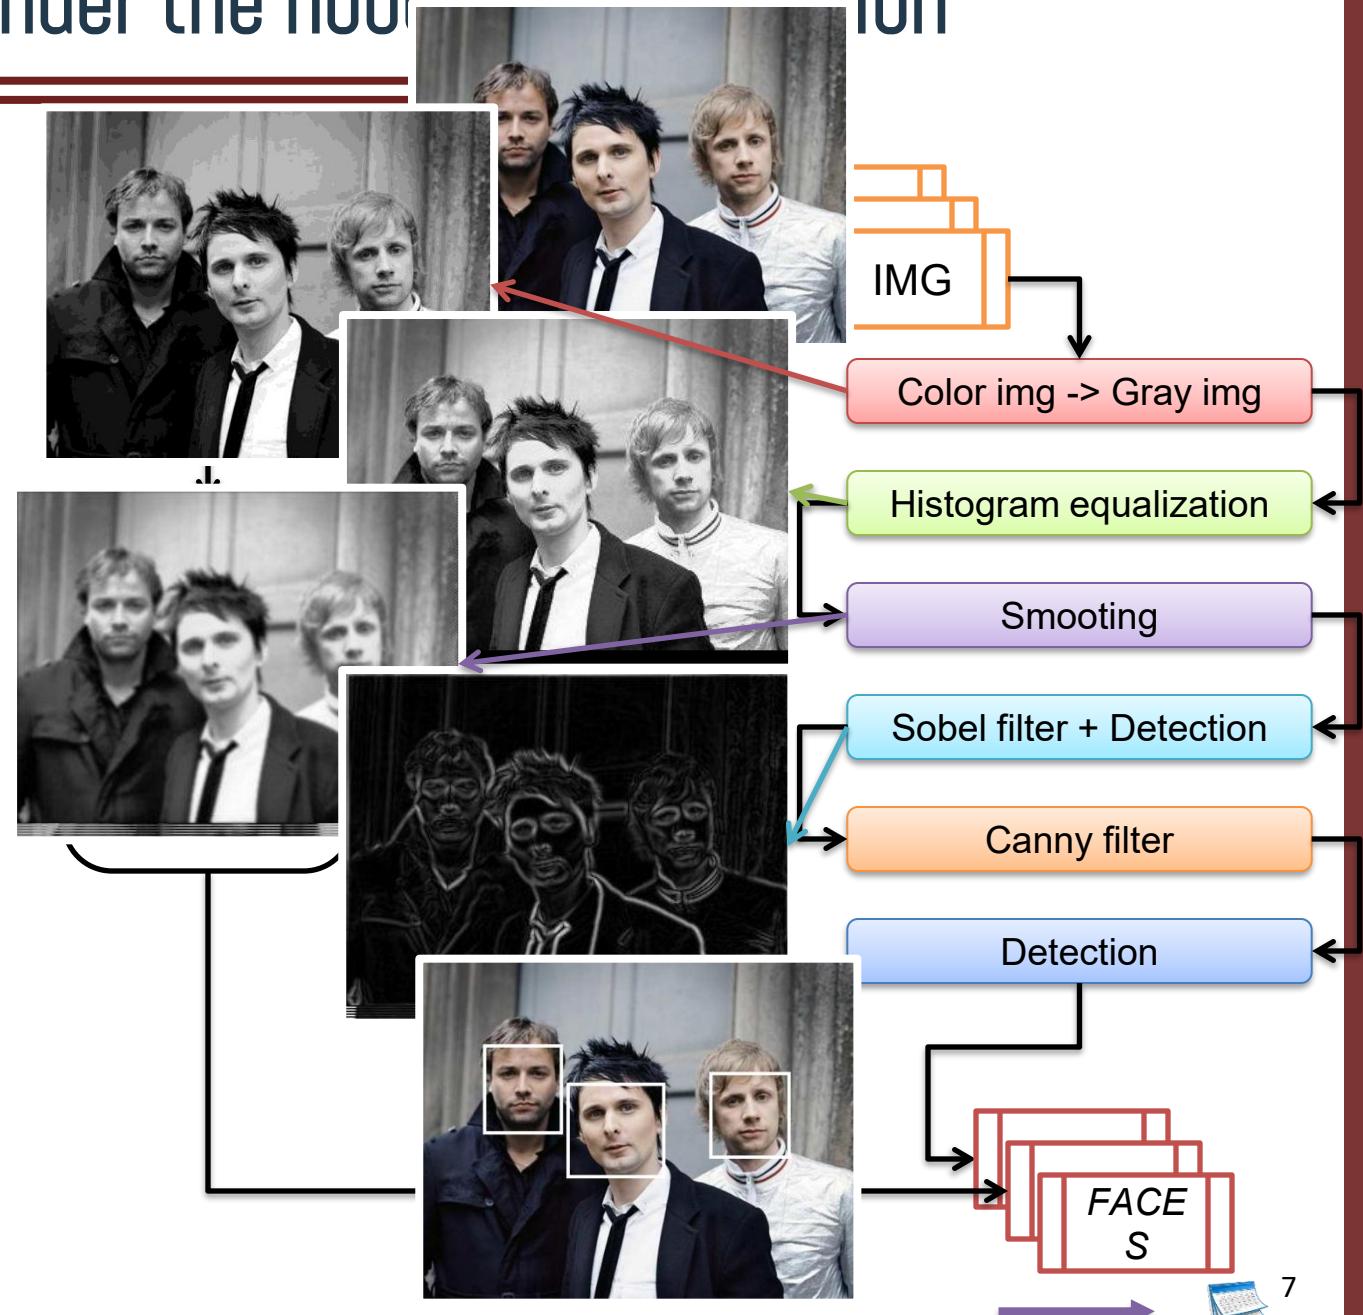




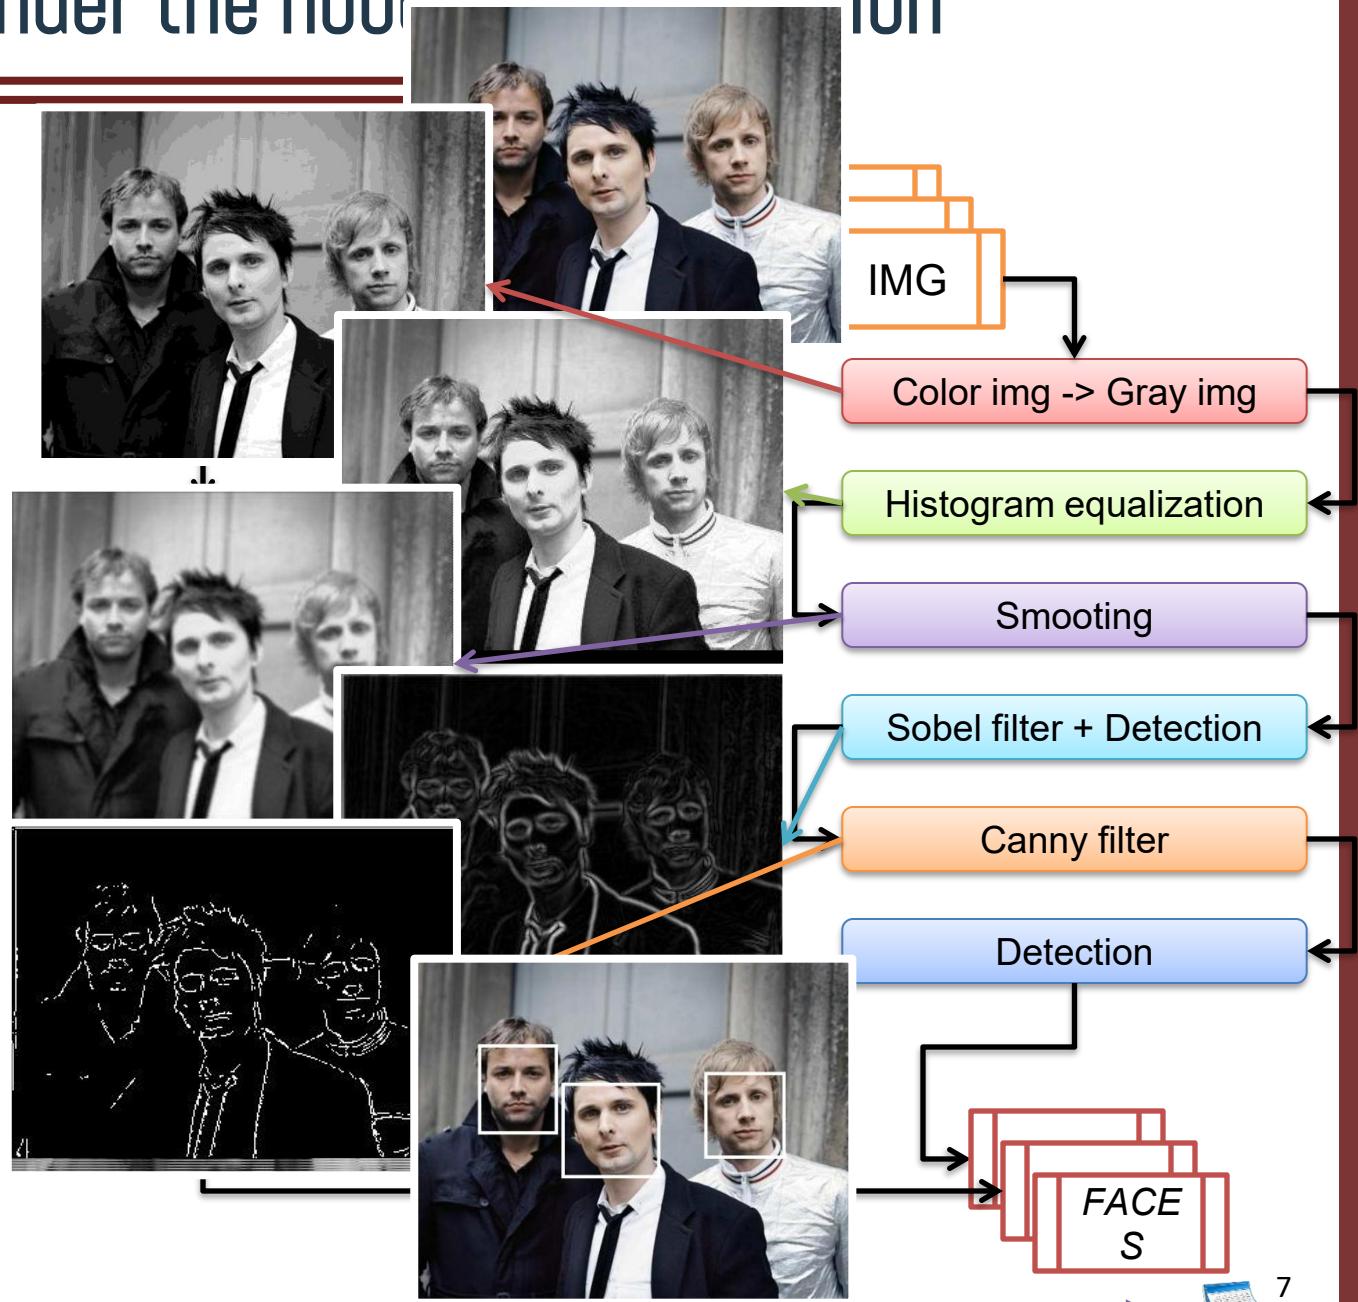
# Under the hood: face detection



# Under the hood: face detection



# Under the hood: face detection



# Image binarization

- › Graylevel image => B/W image
- › Pixel: 256 shades of gray
  - unsigned chars
  - 255 => white
  - 0 => black



```
#define GRAY_THRESHOLD 100
#define WHITE 255
#define BLACK 0
void binarizeImage(const unsigned char inputImg[],
                  unsigned char outputImg[],
                  unsigned int imgDim)
{
    for(int i=0; i<imgDim; i++)
        if(inputImg[i] >= GRAY_THRESHOLD)
            outputImg[i] = WHITE;
        else
            outputImg[i] = BLACK;
}
```



# Image binarization

- › Graylevel image => B/W image
- › Pixel: 256 shades of gray
  - unsigned chars
  - 255 => white
  - 0 => black



```
#define GRAY_THRESHOLD 100
#define WHITE 255
#define BLACK 0
void binarizeImage(const unsigned char inputImg[],
                  unsigned char outputImg[],
                  unsigned int imgDim)
{
    for(int i=0; i<imgDim; i++)
        if(inputImg[i] >= GRAY_THRESHOLD)
            outputImg[i] = WHITE;
        else
            outputImg[i] = BLACK;
}
```



# Image binarization

- › Graylevel image => B/W image
- › Pixel: 256 shades of gray
  - unsigned chars
  - 255 => white
  - 0 => black



```
#define GRAY_THRESHOLD 100
#define WHITE 255
#define BLACK 0
void binarizeImage(const unsigned char inputImg[],
                   unsigned char outputImg[],
                   unsigned int imgDim)
{
    for(int i=0; i<imgDim; i++)
        if(inputImg[i] >= GRAY_THRESHOLD)
            outputImg[i] = WHITE;
        else
            outputImg[i] = BLACK;
}
```

Multiple Data



# Image binarization

- › Graylevel image => B/W image
- › Pixel: 256 shades of gray
  - unsigned chars
  - 255 => white
  - 0 => black



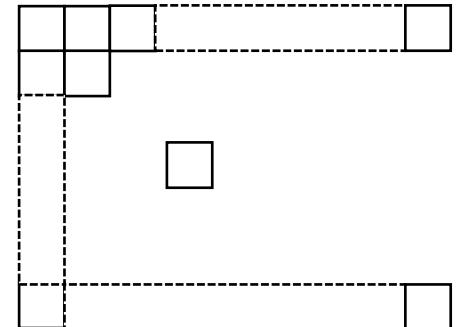
```
#define GRAY_THRESHOLD 100
#define WHITE 255
#define BLACK 0
void binarizeImage(const unsigned char inputImg[],
                   unsigned char outputImg[],
                   unsigned int imgDim)
{
    for(int i=0; i<imgDim; i++)
        if(inputImg[i] >= GRAY_THRESHOLD)
            outputImg[i] = WHITE;
        else
            outputImg[i] = BLACK;
}
```

Single Program



# GPUs

- › Let's (re)design them!
- › We want to perform graphics
  - E.g., filters, shaders...
- › Ultimately, operations on pixels!
  - Same algorithm repeated for each (subset of) pixels
- › Algorithm => program
- › (subset of) pixels => data
- › Same (single) Program, Multiple Data – SPMD
  - Not SIMD!





# A (programmable) machine

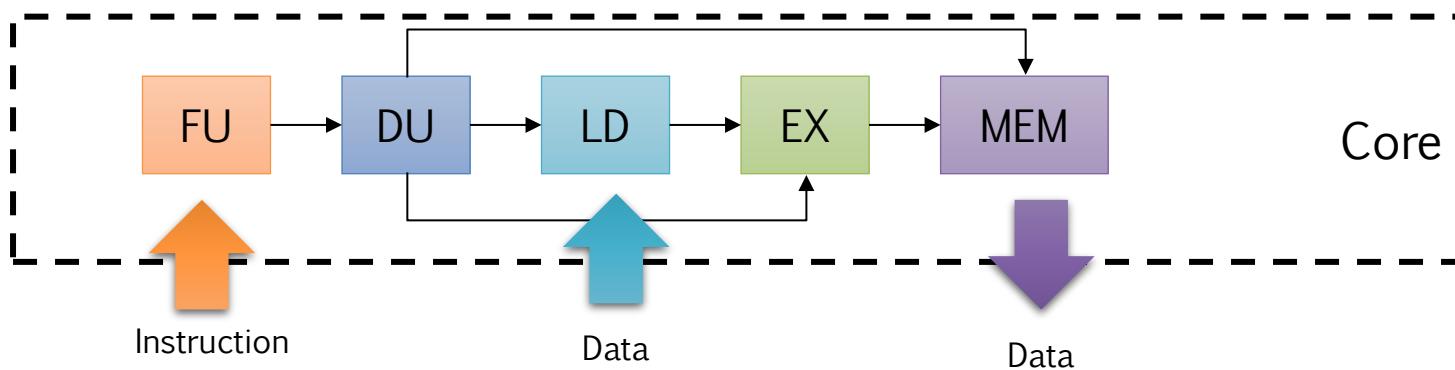
---

---

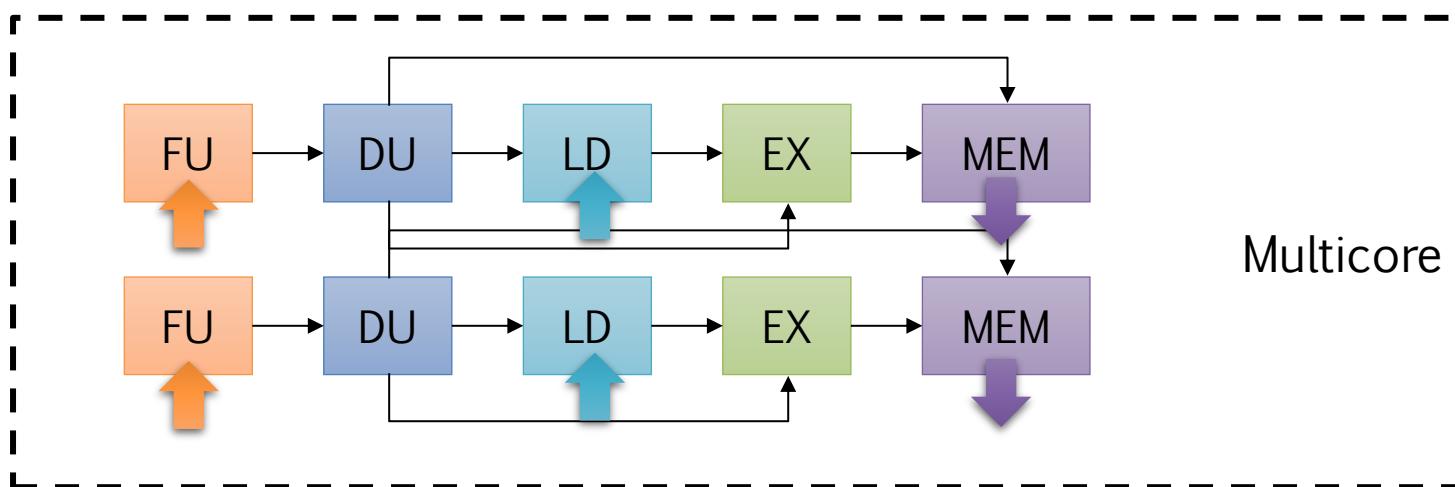
- › Algorithms for image processing are
  - Highly regular (loop-based, with well known boundaries at image rows/columns)
  - Massively parallel (thousands of threads)
- › Regular, "big" loops
  - Single Program (Loop Iteration) Multiple Data - SPMD
  - Parallel threads perform the very same operation on adjacent data
- › We need a massively parallel machine
  - Thousands of cores
- › With simple cores
  - FP Support
- › To perform the very same instruction!
  - Same Fetch Unit and Decode Unit

# Fetch and decode units

- › Traditional pipeline

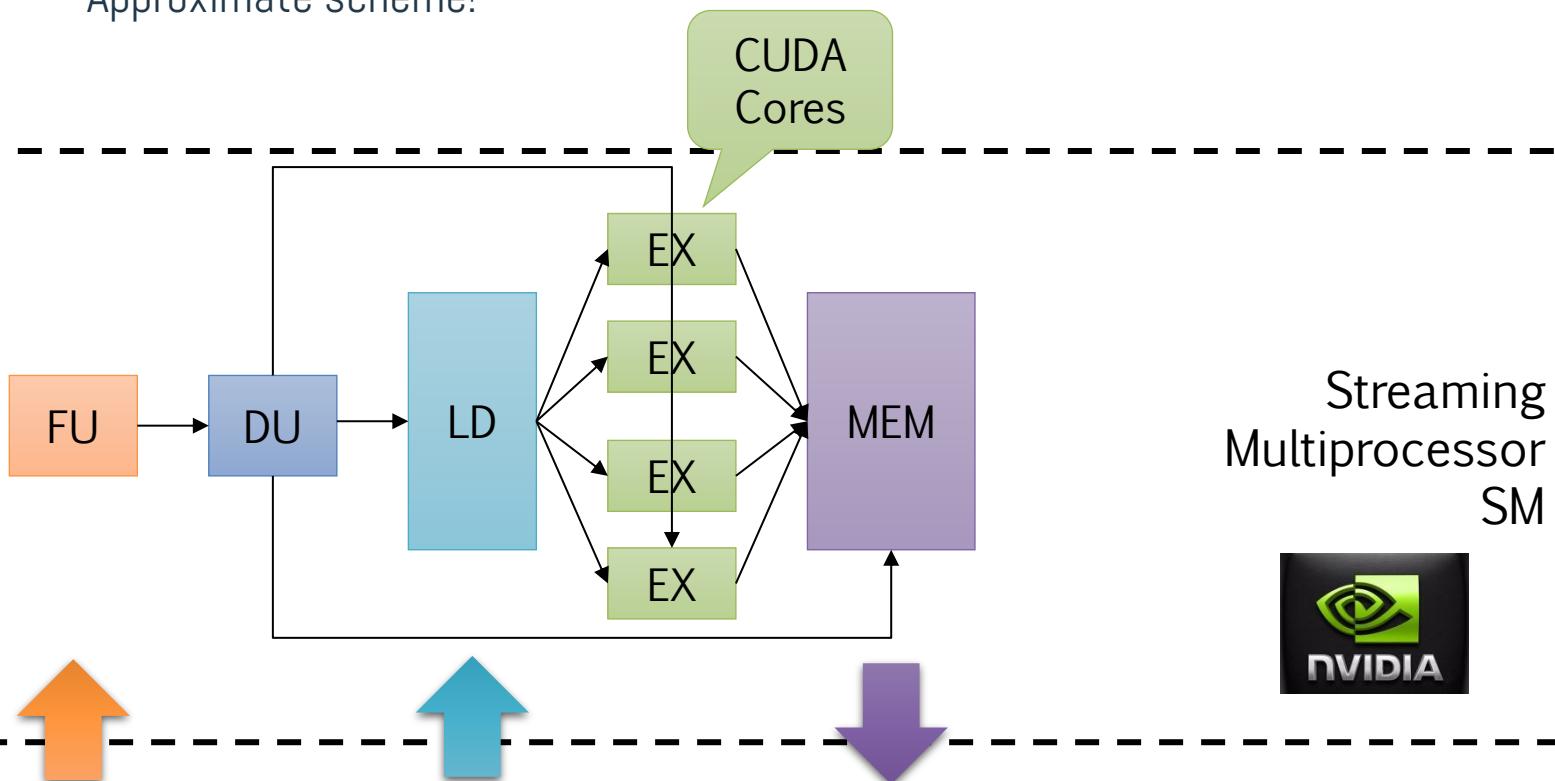


- › Traditional parallel pipeline



# GPU multi-core

- › Share FU, DU, MEM units
  - Approximate scheme!



# SMs as building block

## › Architecture of the SM

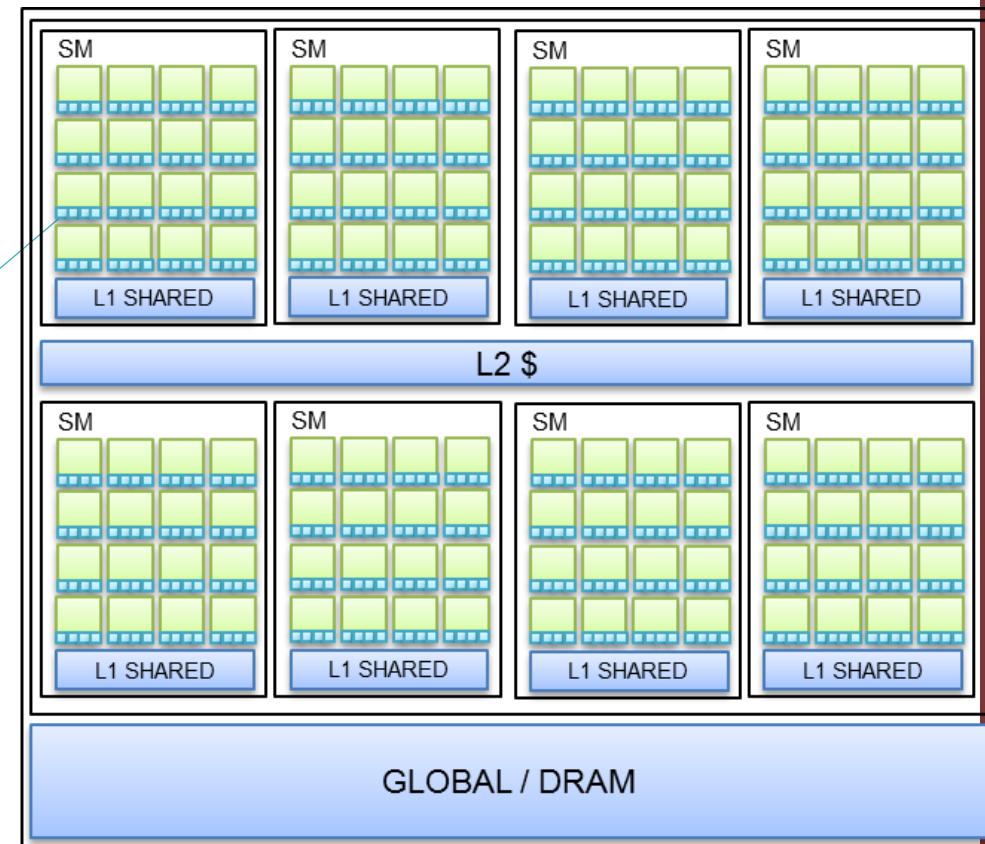
- GPU "class"
- Kepler has 192 cores
- Maxwell/Pascal has 128 cores

## › Number of SMs

- GPU model
- Maxwell's GTX980 has 10
- Pascal's GTX1080 has 20
- Pascal's Drive PX1 has 2

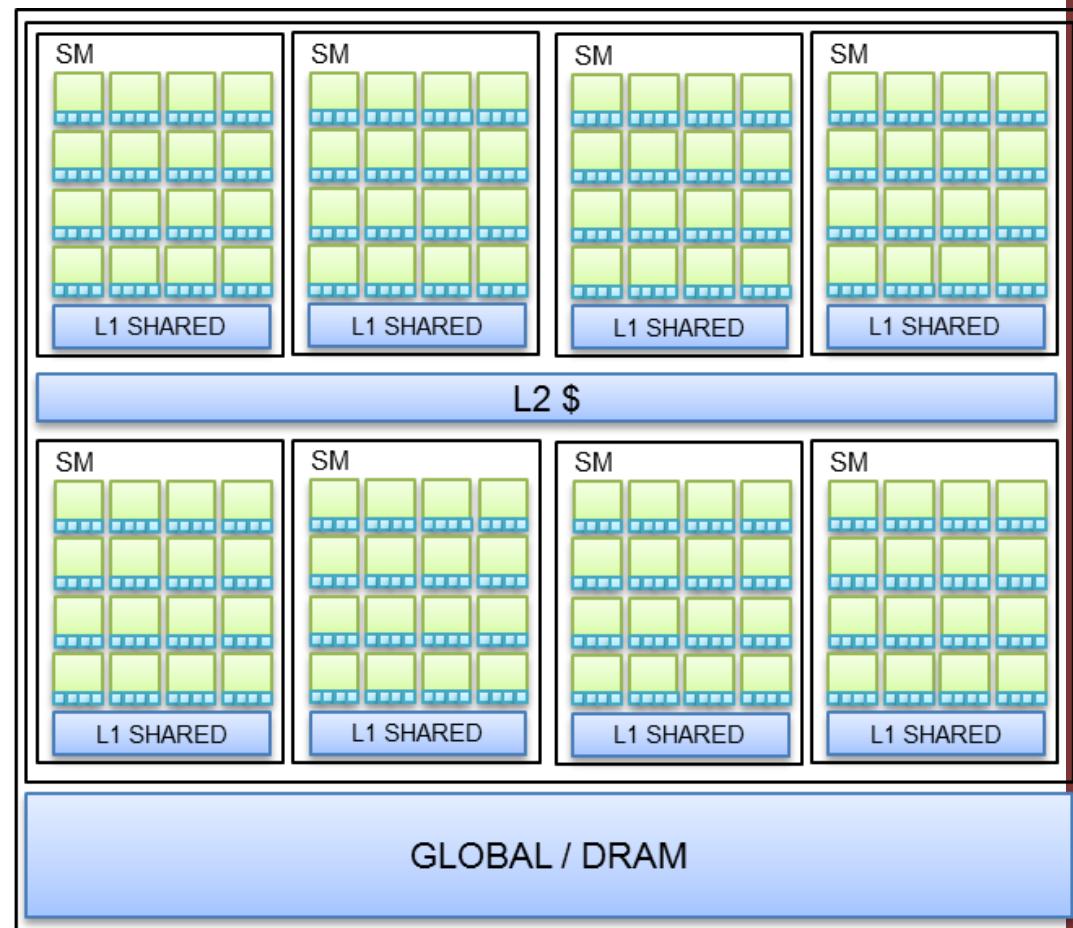
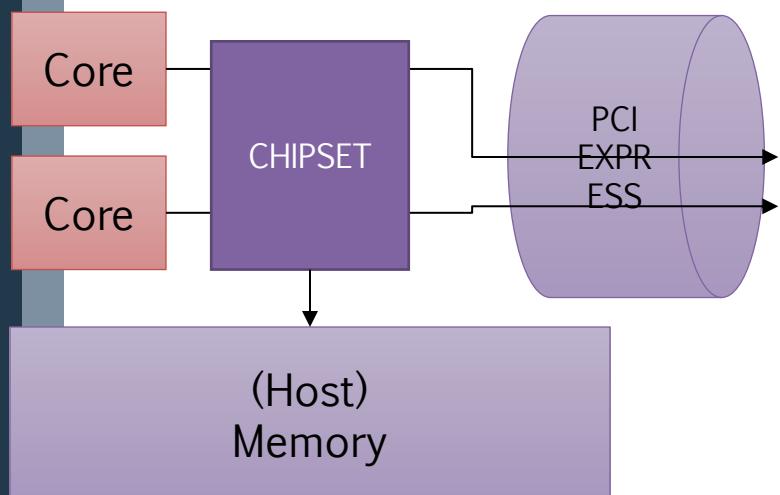
## › NUMA memory system

Local  
Memory



# GPU as a device (Discrete GPGPUs)

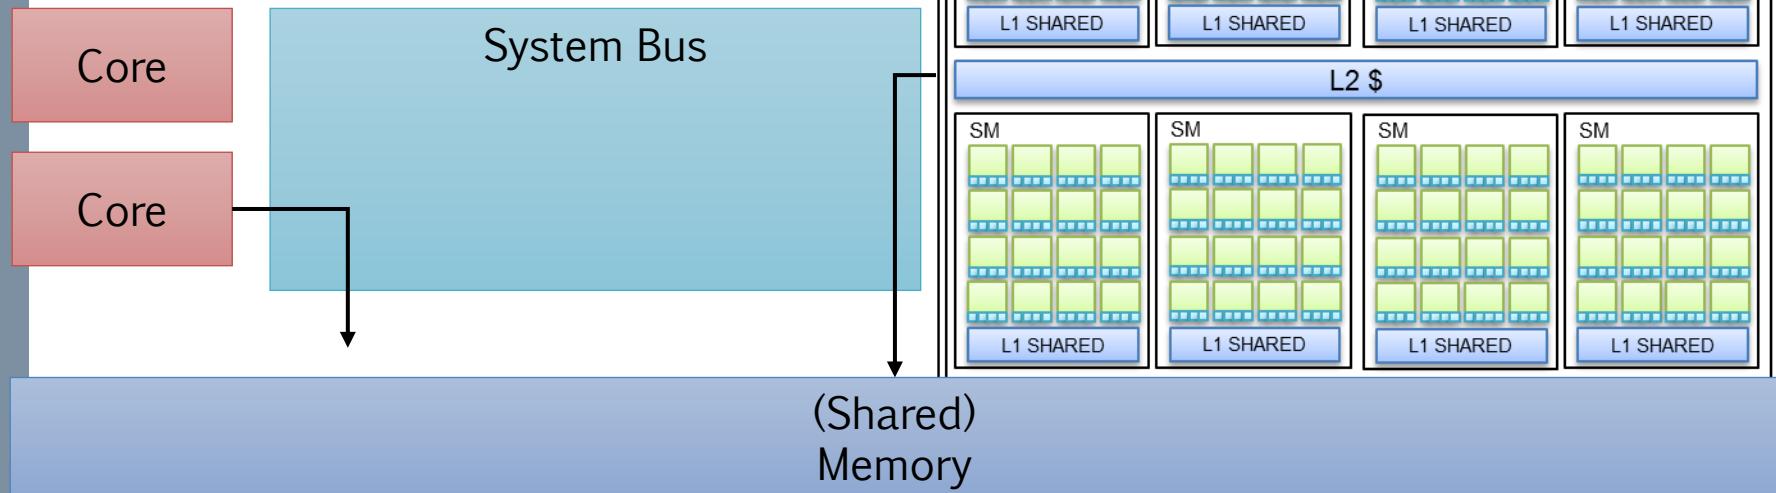
- › Host-device scheme
- › Hierarchical NUMA space
  - Non-Uniform Mem Access



# Integrated GPGPUs

GP-GPU based embedded platforms

- › As opposite to, traditional "discrete" GP-GPUs
- › Still, host + accelerator model
- › Communicate via shared memory
  - No PCI-express
  - CUDA "Unified Virtual Memory"





# To summarize...

---

---

- › Tightly-coupled SMs
  - Multiple cores sharing HW resources: L1 cache, Fetch+Decode Unit, (maybe even) Memory controllers, DSPs...
  - GPU "Class" (NVIDIA Kepler, Maxwell, Pascal, Volta..)
  - ~100s cores
- › Multiple SMs integrated onto one chip
  - GPU "name" (NVIDIA GTX980, GT640...)
  - 1000s cores
  - NUMA hierarchy
- › Typically (but not only) used as co-processor/accelerator
  - PCIEPRESS connectivity
  - Shared memory



# (GP)GPU programming stack

Application(s)

OpenGL



HW



Application(s)

Host

Device

OS  
hooks  
(APIs)

CUDA/OpenCL

API

API

Runtime

Runtime

OS

PCI

???

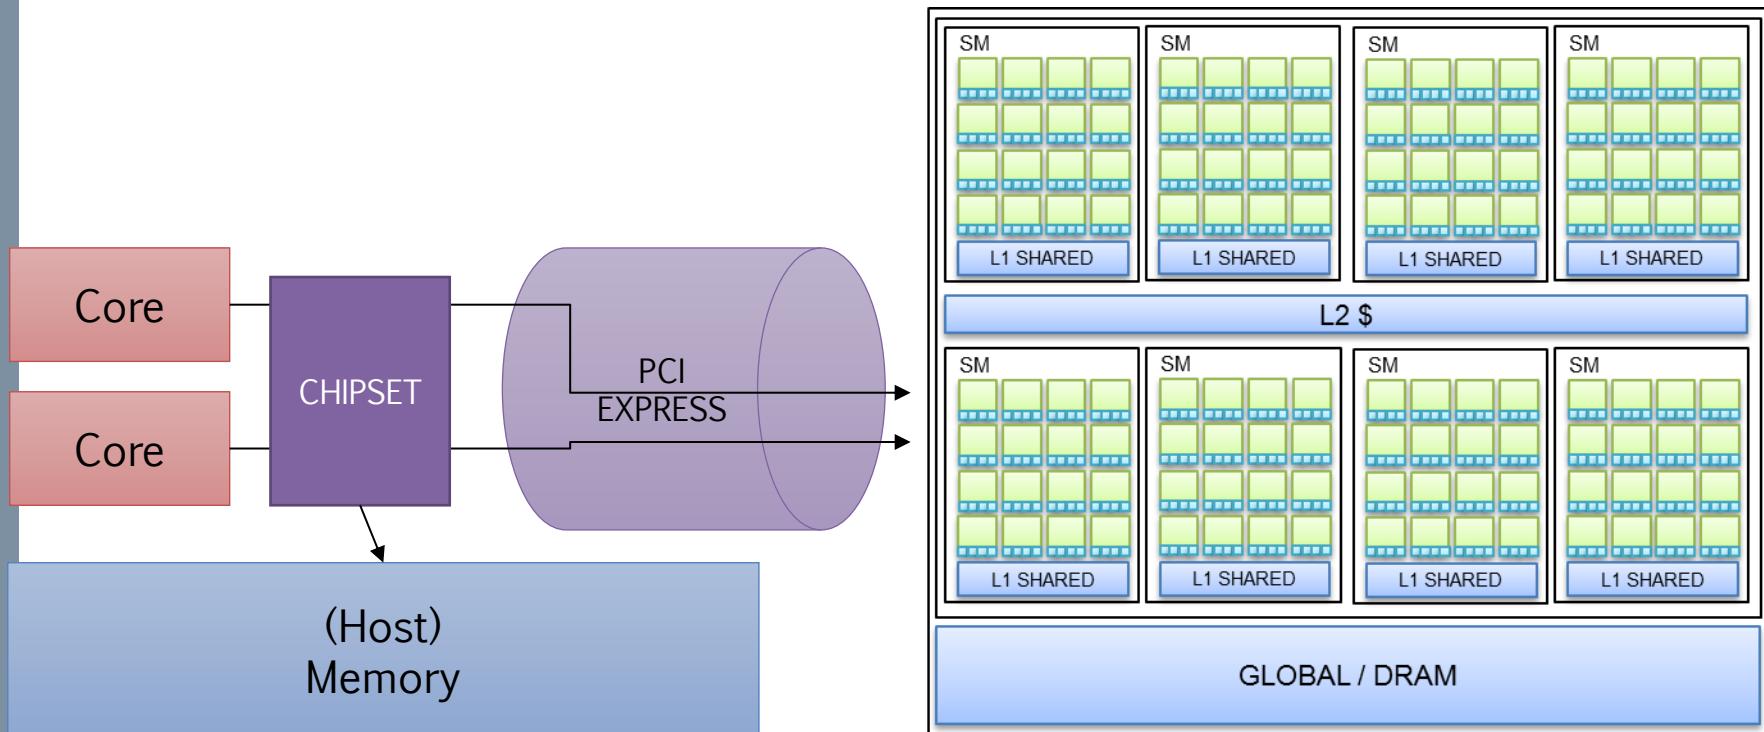


HW



# GPU programming

- › We need a programming model that provides
  1. Simple offloading subroutines
  2. An easy way to write code which runs on thousand threads
  3. A way to exploit the NUMA hierarchy

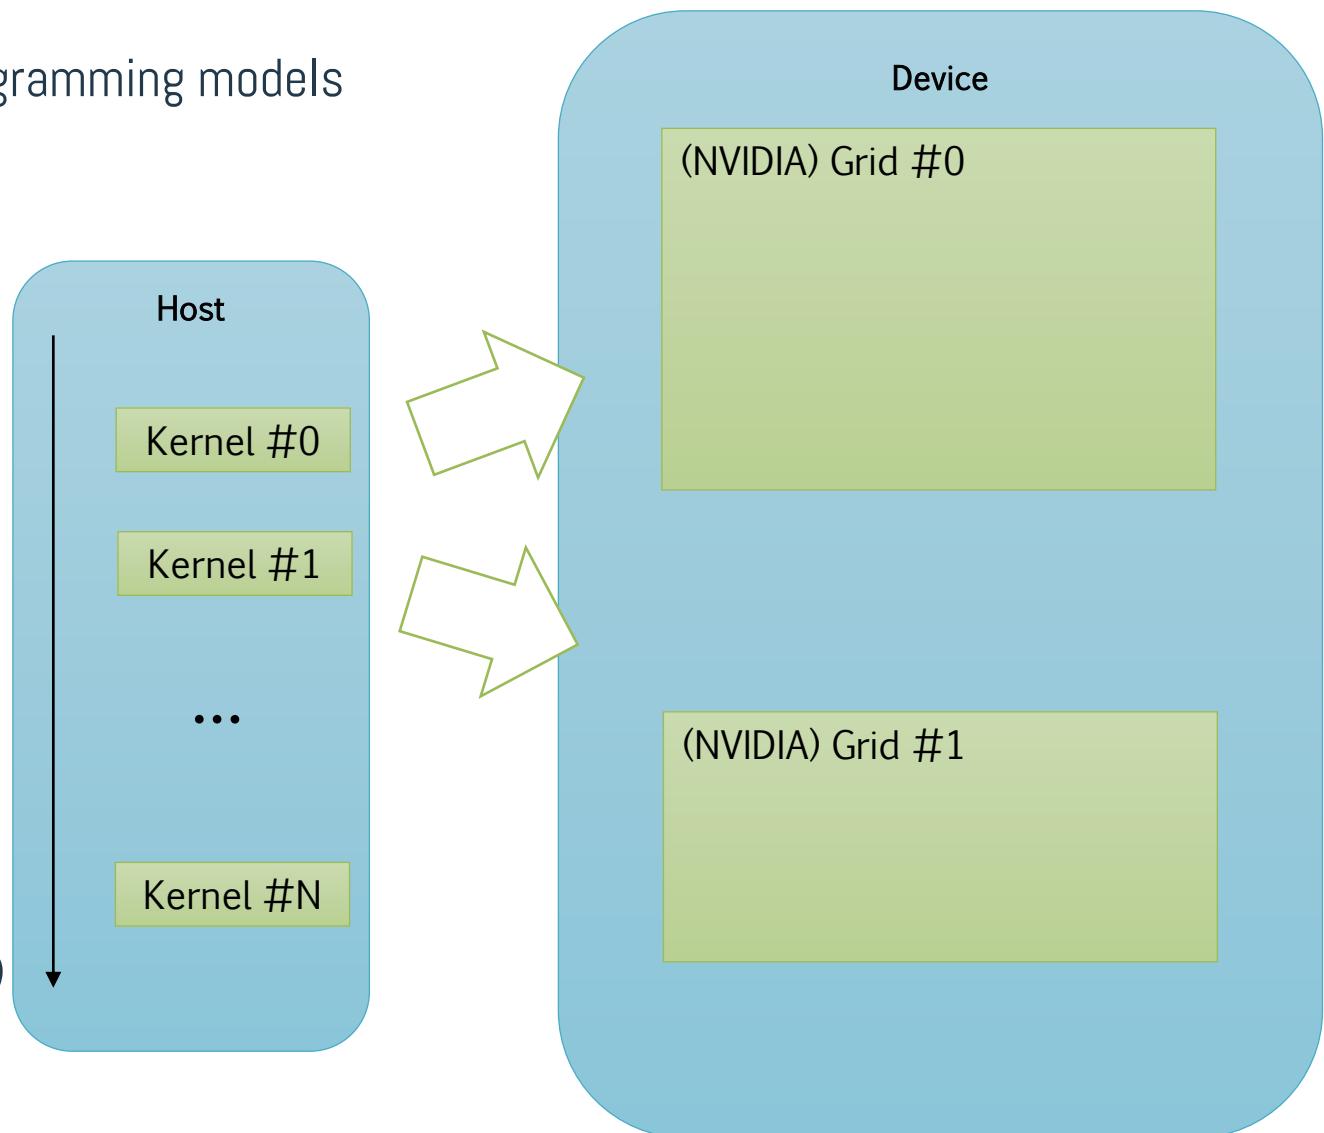


# 1) Offload-based programming

Offload-based programming models

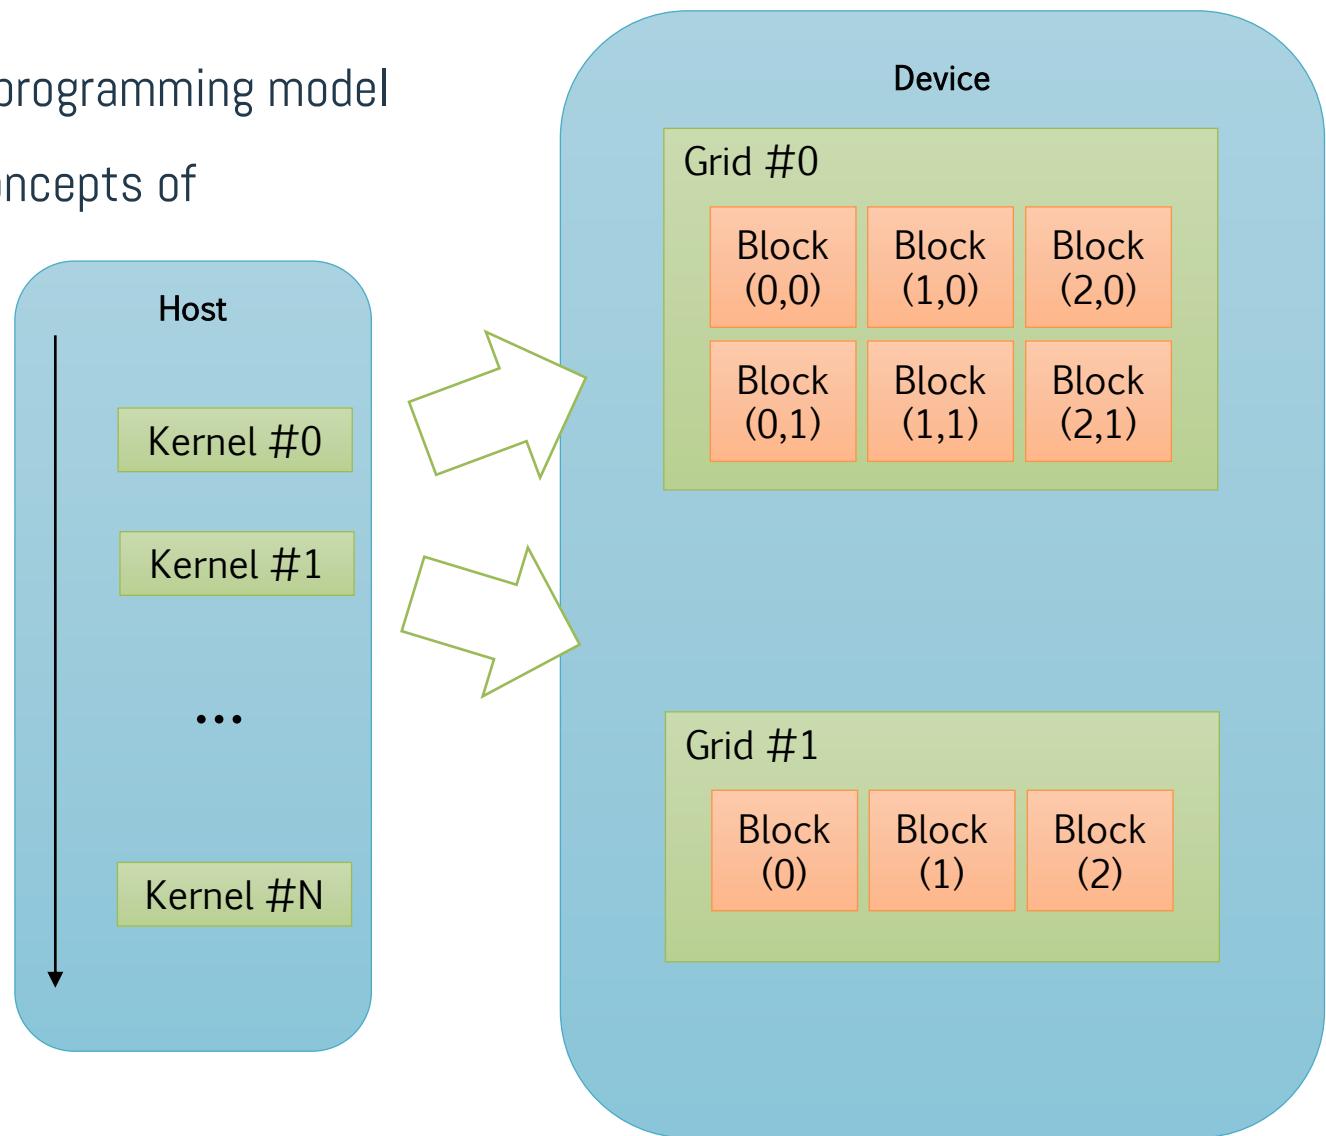
- › CUDA
- › OpenCL
- › OpenMP 4.5

*Portions of  
Code to offload  
⇒ KERNELS  
(typically, loops)*



## 2) Parallelism in CUDA

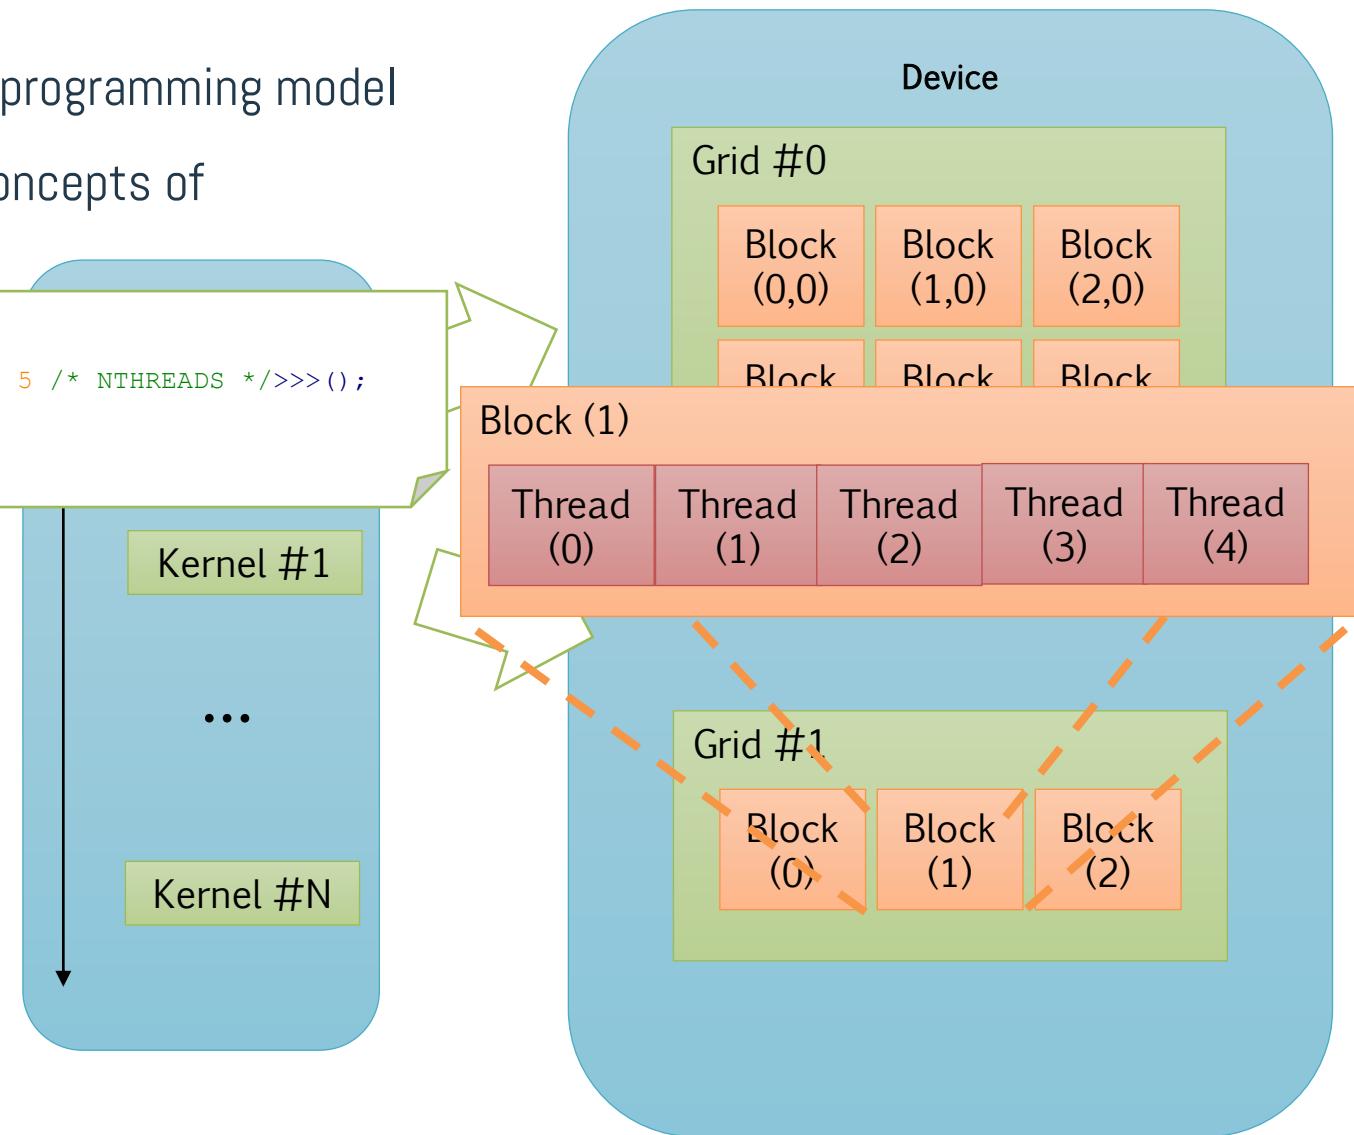
- › Exposed in the programming model
- › Based on the concepts of
  - Grid(s)
  - Block(s)
  - Thread(s)



## 2) Parallelism in CUDA

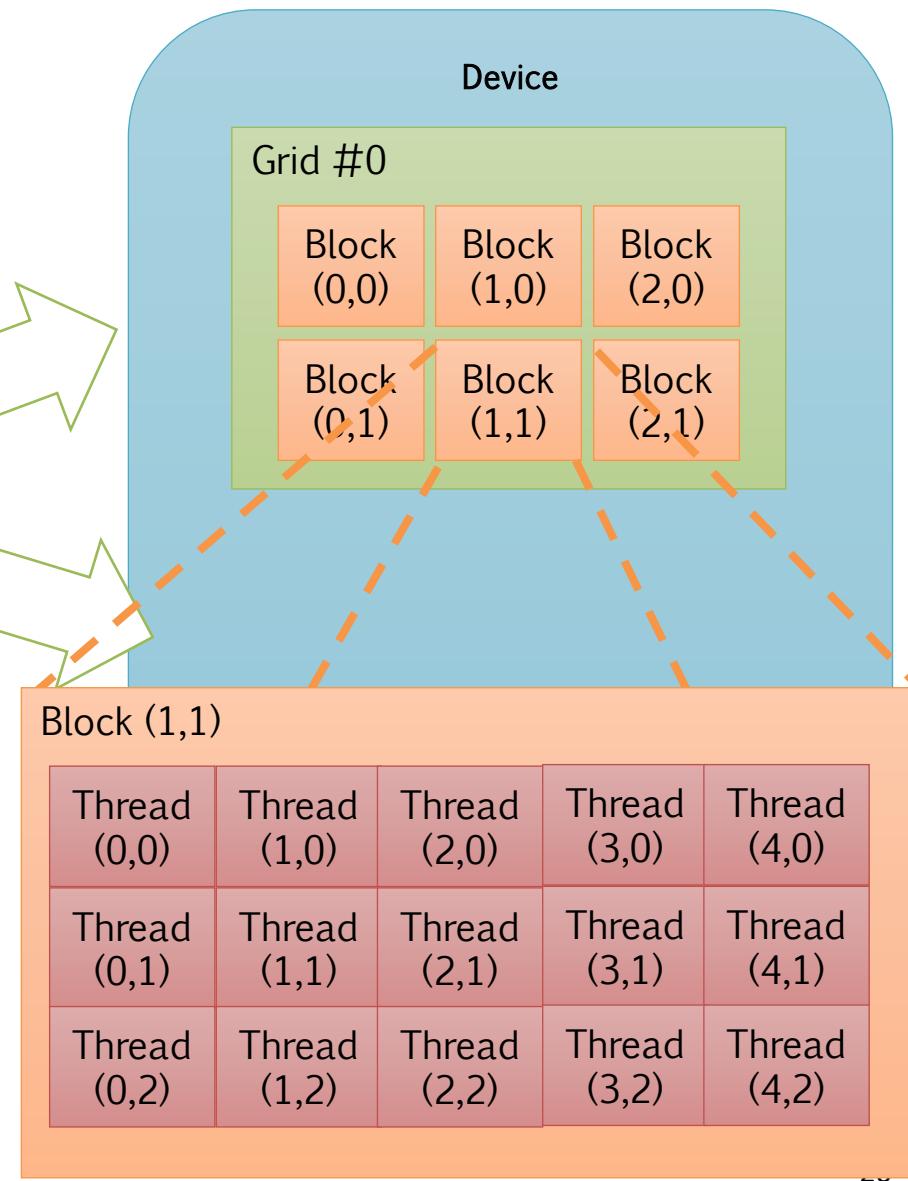
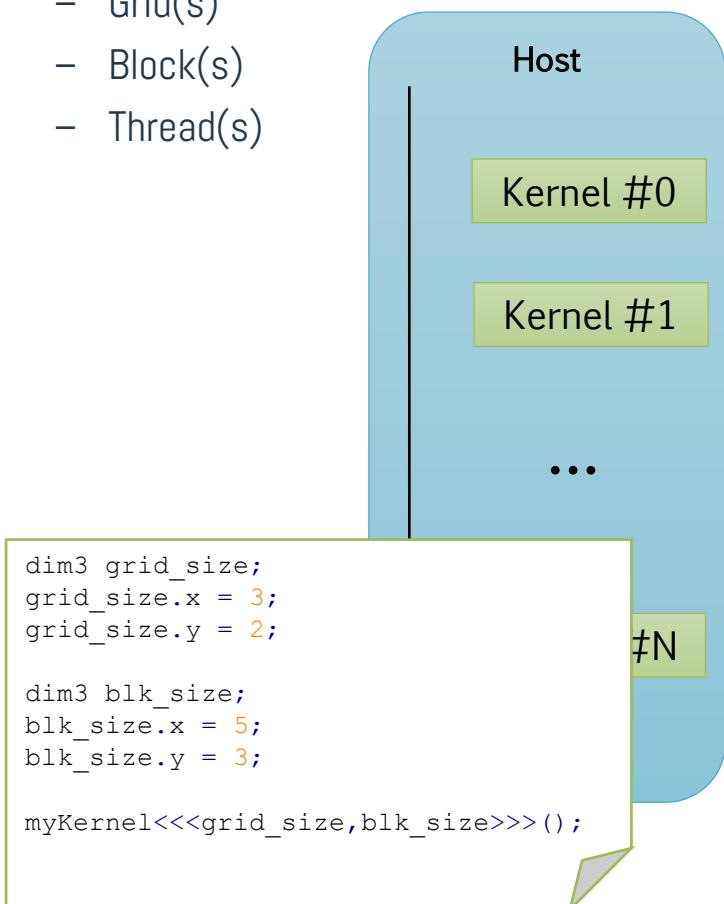
- › Exposed in the programming model
- › Based on the concepts of
  - Grid(s)

```
myKernel<<<3 /* NBLOCKS */, 5 /* NTHREADS */>>>();
```



## 2) Parallelism in CUDA

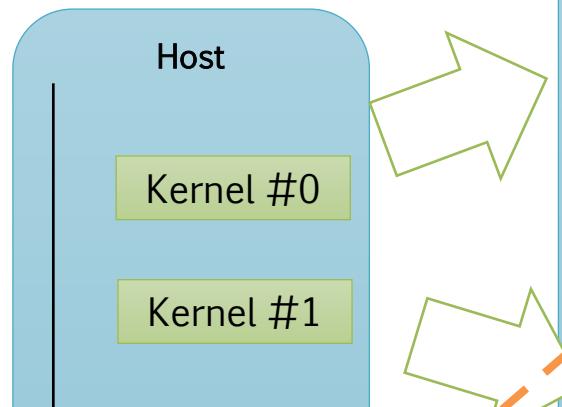
- › Exposed in the programming model
- › Based on the concepts of
  - Grid(s)
  - Block(s)
  - Thread(s)



## 2) Parallelism in CUDA

- › Exposed in the programming model
- › Based on the concepts of

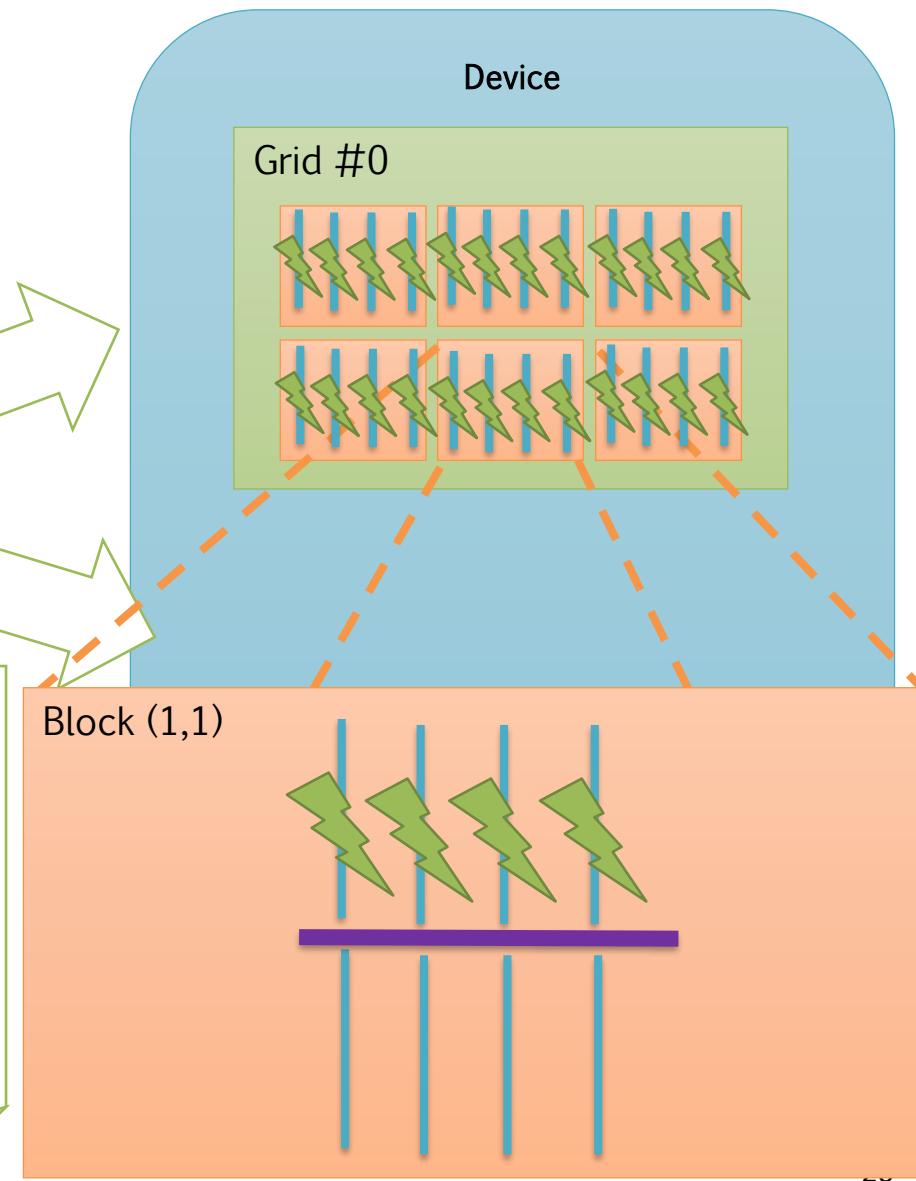
- Grid(s)
- Block(s)
- Thread(s)



```
__global__ void myKernel(
    int *c, const int *a, const int *b)
{
    // Some work

    // Synchronization point for
    // all threads of the same Block
    __syncthreads();

    // Some work
}
```





# Complexity of GPUs

---

---

- › Grids → kernels
- › Blocks X Threads represent a "work-space"
  - Synchronization is possible only within the same CUDA Block
    - › `__syncthreads()`
  - Each thread retrieves its "point" inside this space, and maps it on a specific
    - › Data item, such as array element, matrix element, matrix row...
    - › "Job item", such as a function
    - › Can be 1x1D, 2x2D, 3x3D: extremely (too much) flexible and scalable

# Complexity of GPUs

- › Grids → kernels
- › Blocks X Threads represent a "work-space"
  - Synchronization is possible only within the same CUDA Block
    - › `__syncthreads()`
  - Each thread retrieves its "point" inside this space, and maps it on a specific
    - › Data item, such as array element, matrix element, matrix row...
    - › "Job item", such as a function

```
#define GRAY_THRESHOLD 100
#define WHITE 255
#define BLACK 0
void binarizeImage(const unsigned char inputImg[],
                    unsigned char outputImg[],
                    unsigned int imgDim)
{
    for(int i=0; i<imgDim; i++)
        if(inputImg[i] >= GRAY_THRESHOLD)
            outputImg[i] = WHITE;
        else
            outputImg[i] = BLACK;
}
```

(too much) flexible and scalable



# Complexity of GPUs

- › Grids → kernels
- › Blocks X Threads represent a "work-space"
  - Synchronization is possible only within the same CUDA Block
    - › `__syncthreads()`
  - Each thread retrieves its "point" inside this space, and maps it on a specific
    - › Data item, such as array element, matrix element, matrix row...
    - › "Job item", such as a function

```
#define GRAY_THRESHOLD 100
#define WHITE 255
#define BLACK 0
void binarizeImage(const unsigned char inputImg[],
                    unsigned char outputImg[],
                    unsigned int imgDim)
{
    for(int i=0; i<imgDim; i++)
        if(inputImg[i] >= GRAY_THRESHOLD)
            outputImg[i] = WHITE;
        else
            outputImg[i] = BLACK;
}
```



(too much) flexible and scalable

```
/* ... */

// 1 => # Blocks
// imgDim => #Threads
// 1 thread works on each pixel
int thrId = threadIdx.x;
if(inputImg[thrId] >= GRAY_THRESHOLD)
    outputImg[thrId] = WHITE;
else
    outputImg[thrId] = BLACK;

/* ... */
```



# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**

```
/* ... */  
  
// 1 => # Blocks  
// imgDim => #Threads  
// 1 thread works on each pixel  
int thrId = threadIdx.x;  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
    outputImg[thrId] = WHITE;  
else  
    outputImg[thrId] = BLACK;  
  
/* ... */
```



# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**

```
GRAY_THRESHOLD = 150
inputImg[0] = 200
inputImg[1] = 100
```

```
/* ... */

// 1 => # Blocks
// imgDim => #Threads
// 1 thread works on each pixel
int thrId = threadIdx.x;
if(inputImg[thrId] >= GRAY_THRESHOLD)
    outputImg[thrId] = WHITE;
else
    outputImg[thrId] = BLACK;

/* ... */
```

# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**

```
GRAY_THRESHOLD = 150
inputImg[0] = 200
inputImg[1] = 100
```



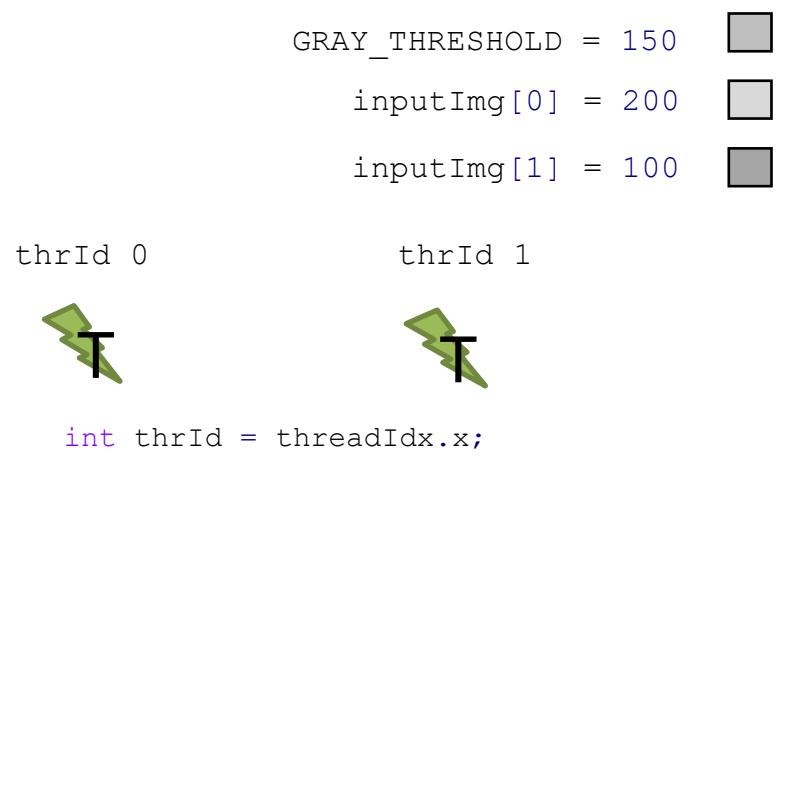
```
/* ... */

// 1 => # Blocks
// imgDim => #Threads
// 1 thread works on each pixel
int thrId = threadIdx.x;
if(inputImg[thrId] >= GRAY_THRESHOLD)
    outputImg[thrId] = WHITE;
else
    outputImg[thrId] = BLACK;

/* ... */
```

# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**



```
/* ... */  
  
// 1 => # Blocks  
// imgDim => #Threads  
// 1 thread works on each pixel  
int thrId = threadIdx.x;  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
    outputImg[thrId] = WHITE;  
else  
    outputImg[thrId] = BLACK;  
  
/* ... */
```



# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
    - Ultimately, the **same Program Counter!!!**
    - Threads cannot do branches - **LOCKSTEP**

```
GRAY_THRESHOLD = 150  
inputImg[0] = 200  
inputImg[1] = 100  
  
thrId 0           thrId 1  
  
int thrId = threadIdx.x;  
  
if(inputImg[thrId] >= GRAY_THRESHOLD)
```

```
/* ... */

// 1 => # Blocks
// imgDim => #Threads
// 1 thread works on each pixel
int thrId = threadIdx.x;
if(inputImg[thrId] >= GRAY_THRESHOLD)
    outputImg[thrId] = WHITE;
else
    outputImg[thrId] = BLACK;

/* ... */
```

# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**



thrId 0



GRAY\_THRESHOLD = 150



inputImg[0] = 200



inputImg[1] = 100



thrId 1



```
int thrId = threadIdx.x;  
  
if(inputImg[thrId] >= GRAY_THRESHOLD)
```

```
/* ... */  
  
// 1 => # Blocks  
// imgDim => #Threads  
// 1 thread works on each pixel  
int thrId = threadIdx.x;  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
    outputImg[thrId] = WHITE;  
else  
    outputImg[thrId] = BLACK;  
  
/* ... */
```

# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**



GRAY\_THRESHOLD = 150



inputImg[0] = 200



inputImg[1] = 100



thrId 0



thrId 1



```
int thrId = threadIdx.x;  
  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
  
outputImg[thrId] = WHITE;
```

```
/* ... */  
  
// 1 => # Blocks  
// imgDim => #Threads  
// 1 thread works on each pixel  
int thrId = threadIdx.x;  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
    outputImg[thrId] = WHITE;  
else  
    outputImg[thrId] = BLACK;  
  
/* ... */
```

# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**



GRAY\_THRESHOLD = 150



inputImg[0] = 200



inputImg[1] = 100



thrId 0



thrId 1



```
int thrId = threadIdx.x;  
  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
  
outputImg[thrId] = WHITE;
```

NOP

```
/* ... */  
  
// 1 => # Blocks  
// imgDim => #Threads  
// 1 thread works on each pixel  
int thrId = threadIdx.x;  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
    outputImg[thrId] = WHITE;  
else  
    outputImg[thrId] = BLACK;  
  
/* ... */
```

# Lockstep

- > (Groups of) cores share the same instruction Fetch/Decode Units
  - Ultimately, the **same Program Counter!!!**
  - Threads cannot do branches - **LOCKSTEP**



GRAY\_THRESHOLD = 150



inputImg[0] = 200



inputImg[1] = 100



thrId 0



thrId 1



```
int thrId = threadIdx.x;  
  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
  
outputImg[thrId] = WHITE;  
  
else
```

NOP

```
/* ... */  
  
// 1 => # Blocks  
// imgDim => #Threads  
// 1 thread works on each pixel  
int thrId = threadIdx.x;  
if(inputImg[thrId] >= GRAY_THRESHOLD)  
    outputImg[thrId] = WHITE;  
else  
    outputImg[thrId] = BLACK;  
  
/* ... */
```



# Lockstep

- › (Groups of) cores share the same instruction Fetch/Decode Units
    - Ultimately, the **same Program Counter!!!**
    - Threads cannot do branches - **LOCKSTEP**



```
GRAY_THRESHOLD = 150  
inputImg[0] = 200  
inputImg[1] = 100
```

thrId 0



NOE



A photograph of Chuck Norris giving a thumbs-up gesture. He has his signature reddish-brown hair and beard. A small Pepsi can is visible in the bottom left corner.

```
/* ... */

// 1 => # Blocks
// imgDim => #Threads
// 1 thread works on each pixel
int thrId = threadIdx.x;
if(inputImg[thrId] >= GRAY_THRESHOLD)
    outputImg[thrId] = WHITE;
else
    outputImg[thrId] = BLACK;

/* ... */
```



# Warps, and lockstep

---

---

- › Threads are grouped in **warps**
  - 1 warp <-> 32 CUDA threads
  - Units of scheduling
  - Threads of a single blocks are scheduled and de-scheduled 32 by 32
- › Threads within the same warp run in **LOCKSTEP**
- › Memory accesses within the single warp are **coalesced**

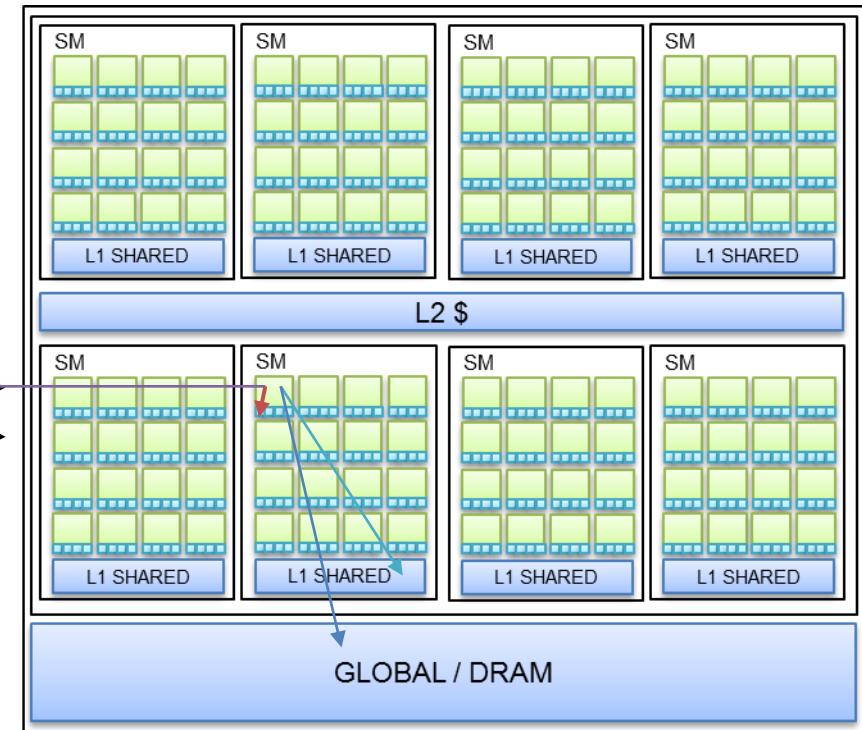
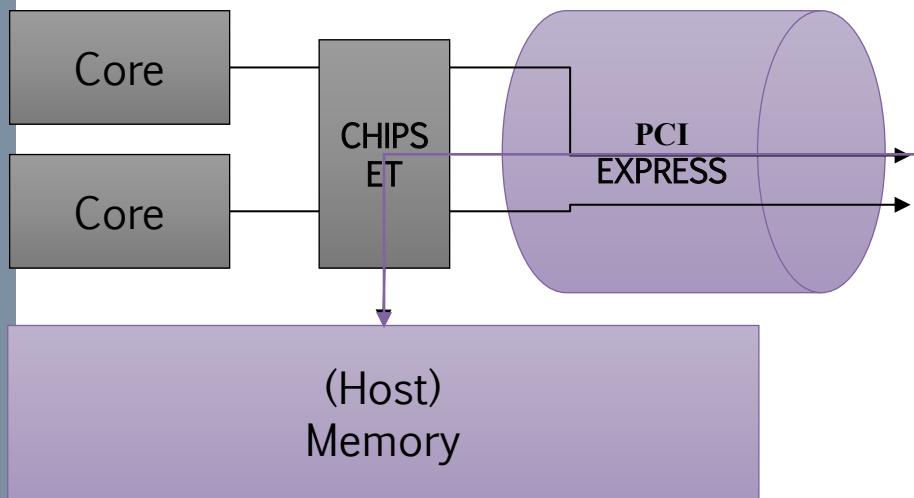
# 3) Exploit NUMA in CUDA

## › Four memory spaces

- Host
- Device **Global**
- Device **Shared**
- Device **Local**

## › Need a way to

- Allocate memory in them
- Move data across them

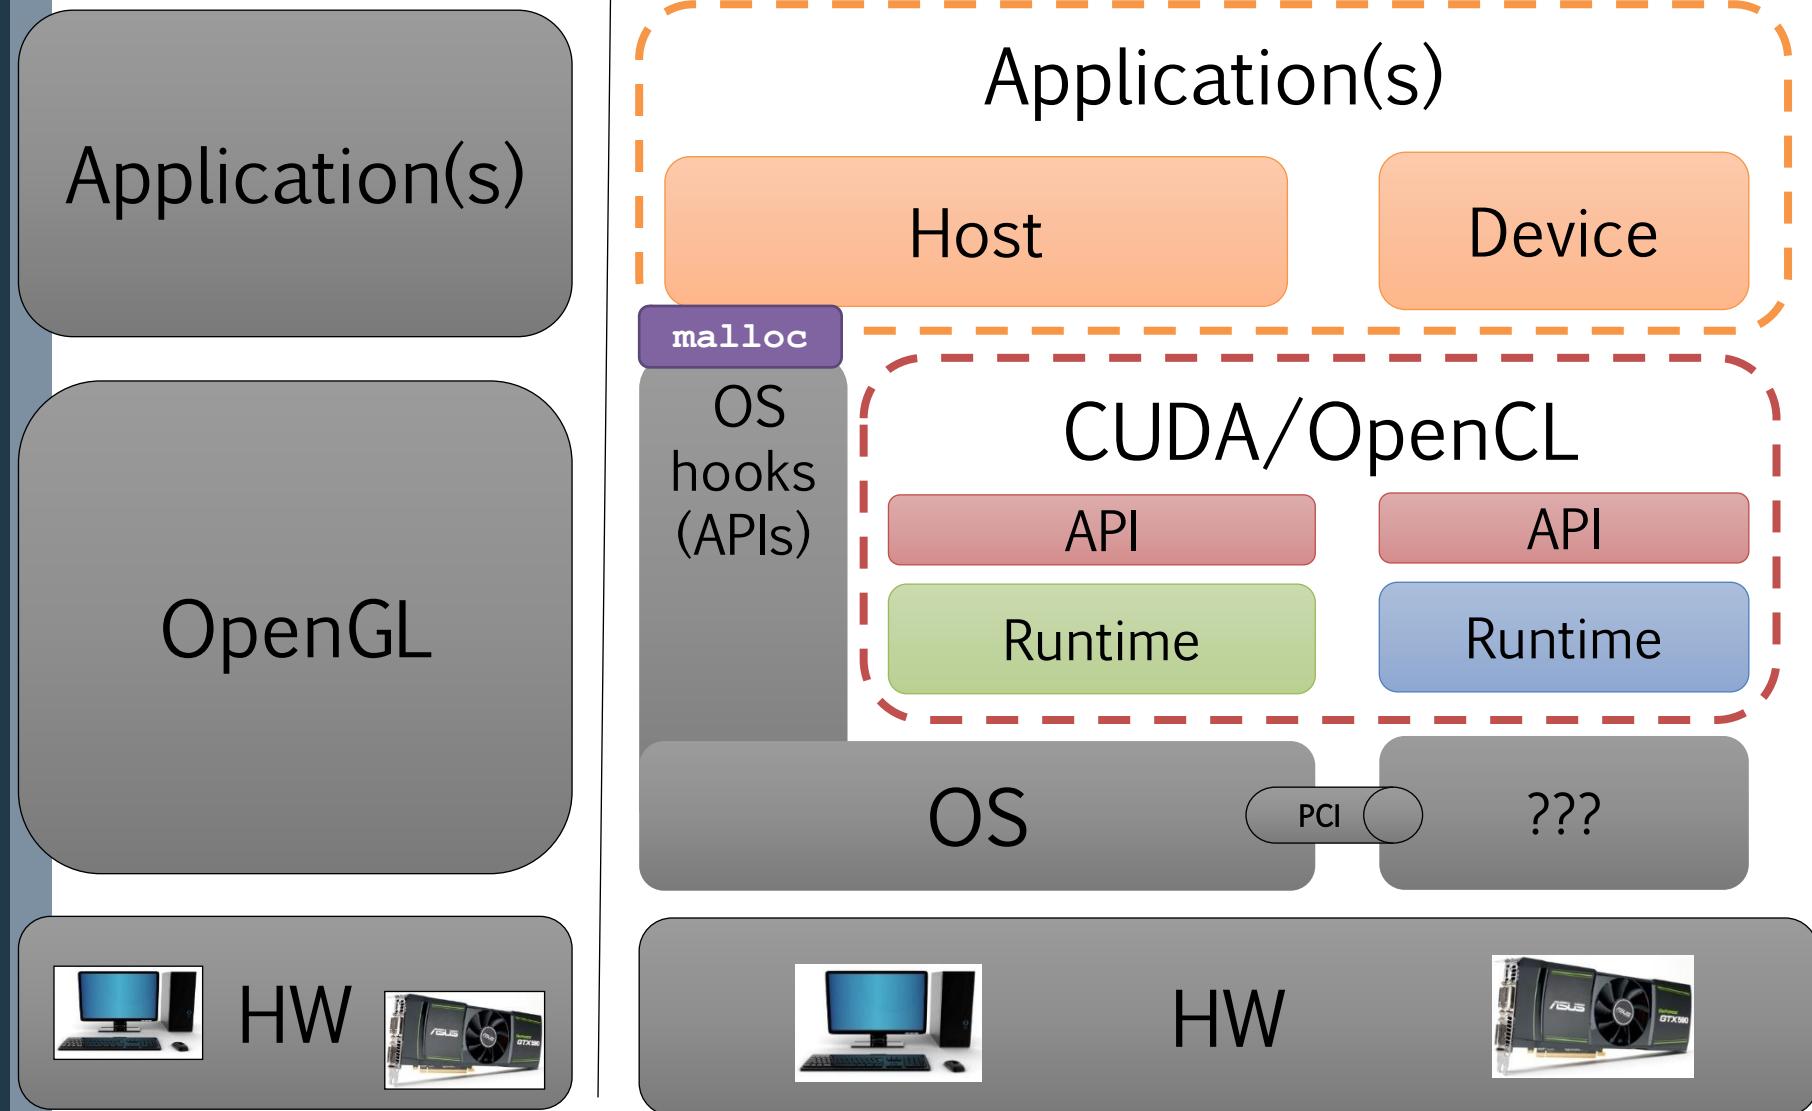


# GPU memory size

	GeForce GT 640 : Liu	GeForce GTX 980 : Turing
<b>Microarchitettura</b>	Kepler	Maxwell
<b>Versione capacità di calcolo</b>	3.0	5.2
<b>Core CUDA</b>	384	2048
<b>Clock del processore</b>	891 MHz	1126 MHz
<b>Clock grafico</b>	900 MHz	1216 MHz
<b>Global memory</b>	2047 MB	4095 MB
<b>Constant memory</b>	64 KB	64 KB
<b>Shared memory per multiprocessor</b>	48 KB	96 KB
<b>Local memory per thread</b>	512 KB	512 KB
<b>Registri a 32-bit per multiprocessor</b>	32 KB	64 KB
<b>Velocità della memoria</b>	1.8 Gbps	7.0 Gbps
<b>Interfaccia della memoria</b>	128-bit DD3	256-bit GDDR5
<b>Supporto del bus</b>	PCI-E 3.0	PCI-E 3.0

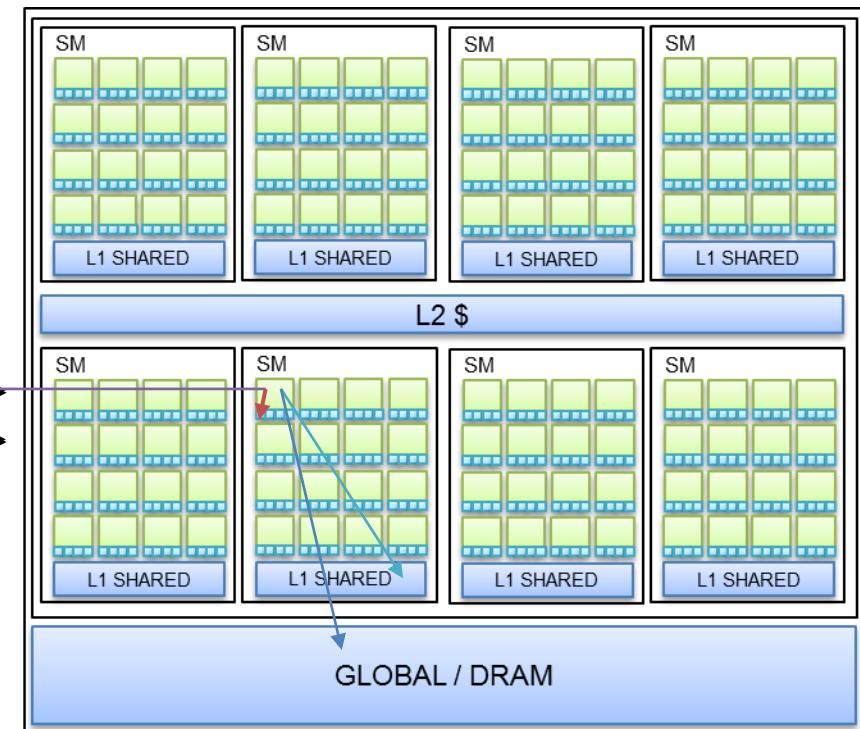
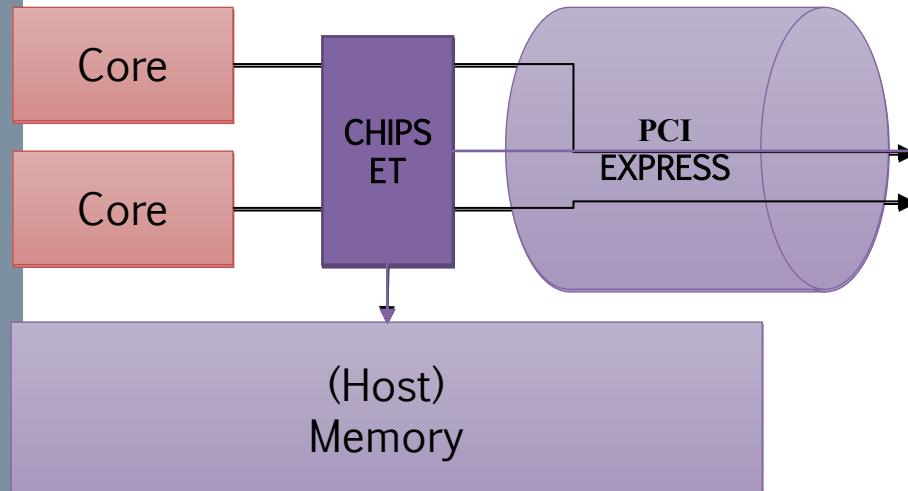


# (GP)GPU programming stack

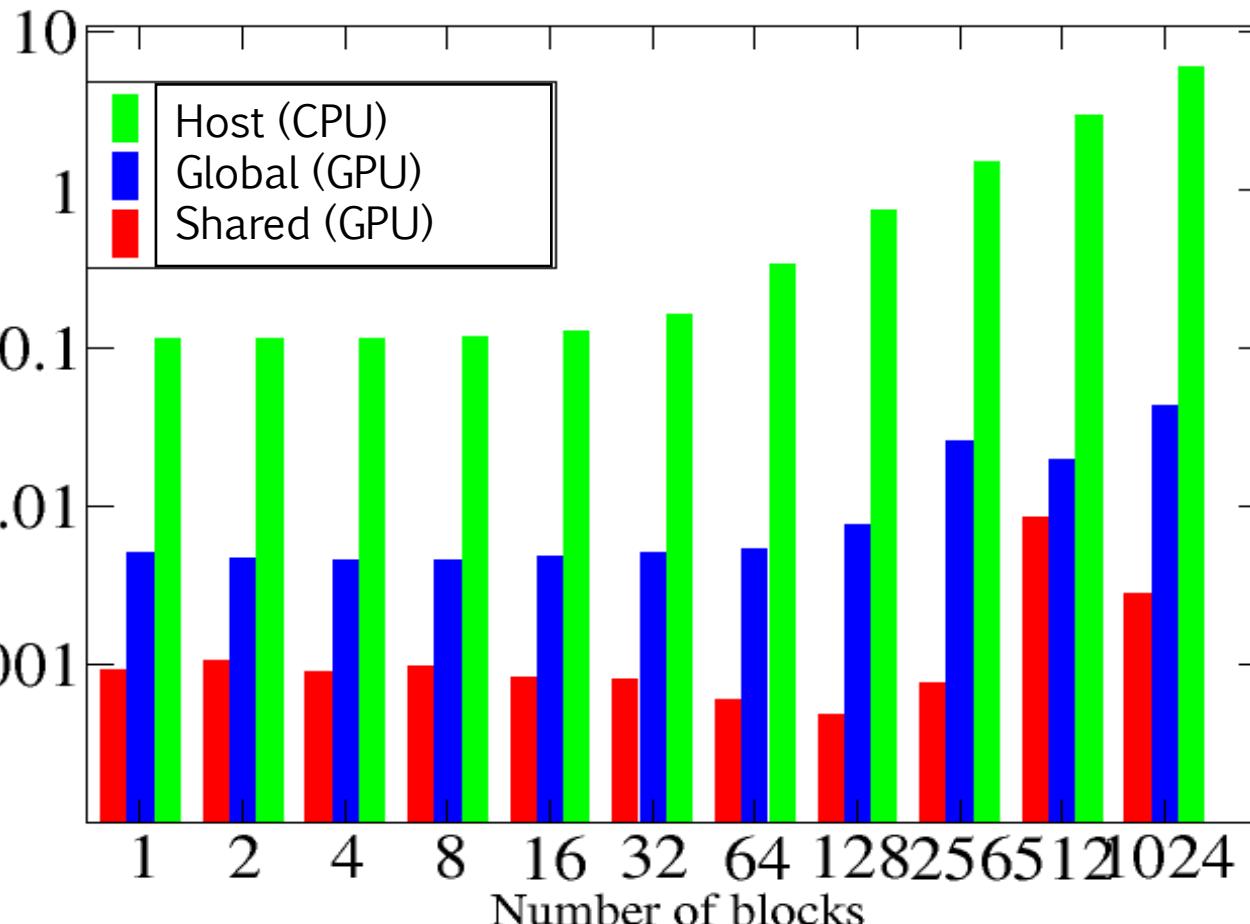


# 3) Exploit NUMA in CUDA

- › Runtime must be aware of all
  - › Memory allocations
    - `cudaHostAlloc` → Host mem
    - `cudaMalloc` → Global mem
    - `__shared__` keyword → Shared mem
  - › Data movements
    - `cudaMemcpy`
    - `cudaMemcpyAsync`



# Non-Uniform Access Time



# OpenCL

- › Open Computing Language
  - More verbose than CUDA
- › More "library-based" approach
- › Different artifacts for managing parallelism
  - CUDA blocks, Threads
  - OpenCL Work Groups, work items

Host



PCI  
EXPRESS

Device





# CUDA vs. OpenCL - Offload code



```
/* Create Command Queue */
command_queue = clCreateCommandQueue(context, device_id, 0, &ret);

/* Create Kernel Program from the source */
program = clCreateProgramWithSource(context, 1, (const char **)&source_str,
                                    (const size_t *) &source_size, &ret);

/* Build Kernel Program */
ret = clBuildProgram(program, 1, &device_id, NULL, NULL, NULL);

/* Create OpenCL Kernel */
kernel = clCreateKernel(program, "hello", &ret);

/* Execute OpenCL Kernel */
ret = clEnqueueTask(command_queue, kernel, 0, NULL,NULL);
```



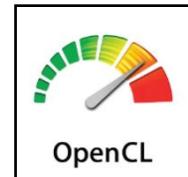
```
helloworld<<<3,5>>>();
cudaDeviceSynchronize()
```



# CUDA vs. OpenCL - Kernel code

```
__kernel void helloworld()
{
    int wiId = get_local_id(0);
    int wgId = get_group_id(0);
    int wiNum = get_local_size(0);
    int wgNum = get_num_groups(0);

    printf("\t\t\t\t[DEVICE] Hello World! \
           I am Work Item #%-d out of %d, \
           and I belong to Work Group #%-d out of %d\n",
           wiId, wiNum, wgId, wgNum);
    return;
}
```



```
__global__ void helloworld()
{
    int thrId = threadIdx.x;
    int blkId = blockIdx.x;
    int thrNum = blockDim.x;
    int blkNum = gridDim.x;

    printf("\t\t\t\t[DEVICE] Hello World! \
           I am thread #%-d out of %d, \
           and I belong to block #%-d out of %d\n",
           thrId, thrNum, blkId, blkNum);

    return;
}
```





# How to run the examples

---

---

Let's  
code!

- › Download the Code/ folder from the course website

- › Compile

```
$ nvcc code.cu [-o myprogram]
```

- › Run

- \$ ./a.out [./myprogram]

# References



## Course website

- › [http://hipert.unimore.it/people/paolob/pub/Industrial\\_Informatics/index.html](http://hipert.unimore.it/people/paolob/pub/Industrial_Informatics/index.html)

## My contacts

- › [pao.lo.burgio@unimore.it](mailto:pao.lo.burgio@unimore.it)
- › <http://hipert.mat.unimore.it/people/paolob/>

## Resources

- › <http://www.nvidia.it/object/cuda-parallel-computing-it.html>
- › <https://www.khronos.org/opencl>
- › A "small blog"
  - <http://www.google.com>